

# Partitioning Tables for Faster Queries

---

# Overview

**Optimize queries with Hive partitions**

**Work with partitioned tables**

- Creating partitions
- Loading data into partitions
- Querying data from partitions

**Implement static and dynamic partitioning on Hive tables**

# Faster Queries on Traditional Databases

---

# Designing Databases



## **Faster queries**

**Retrieving information from a  
database should be fast**

# Designing Databases



**Traditional database admins **design** data layouts based on **usage** patterns**



# Usage Patterns and Design

**Do users access customer and order information together?**

**Customer id is a foreign key in the orders table**

**Create a Customer\_Orders table grouping both information together**




# Usage Patterns and Design

Do users often require order information for a particular date?

Create an **index** on the date column of the orders table


# Indexes



## **Faster lookup**

**Speed up data retrieval in  
traditional relational databases**






# Indexes

```
select * from Orders  
where date = "12/07/2016"
```

```
select * from Orders  
where date > "12/07/2016"  
and date < "16/07/2016"
```




# Indexes

Associated with a table which allow **fast lookup** of records

Created from one or more columns


Copy of column data that can be **searched** very **efficiently**

# Indexes



Well-designed indexes are  
can be a **major performance  
improvement** for queries

# Indexes ~ Partitioning



Partitions are the **logical equivalent** of indexes in Hive

# Partitioning Tables in Hive

---

# Partitioning



**Split data into smaller subsets**


**Separate records into manageable parts based on a column value**

# Partitioning



**Consider an e-commerce site with order data from across the United States**

# Partitioning





# Partitioning




**Records will be split across multiple machines in the cluster**

# Partitioning




**select \* from Orders where state = "WA"**

# Partitioning



**`select * from Orders where state = "WA"`**

# Partitioning



**All records** on  
each machine  
need to be  
scanned to  
retrieve results

`select * from Orders where state = "WA"`

# Partitioning



**select \* from Orders where state = “WA”**

What if the **dataset is huge** on each node?

What if this is the **most common query** run?

# Partitioning

**Data may be naturally split into logical units**



**Customers in the US**

# Partitioning

WA

OR

CA

CT

NY

GA

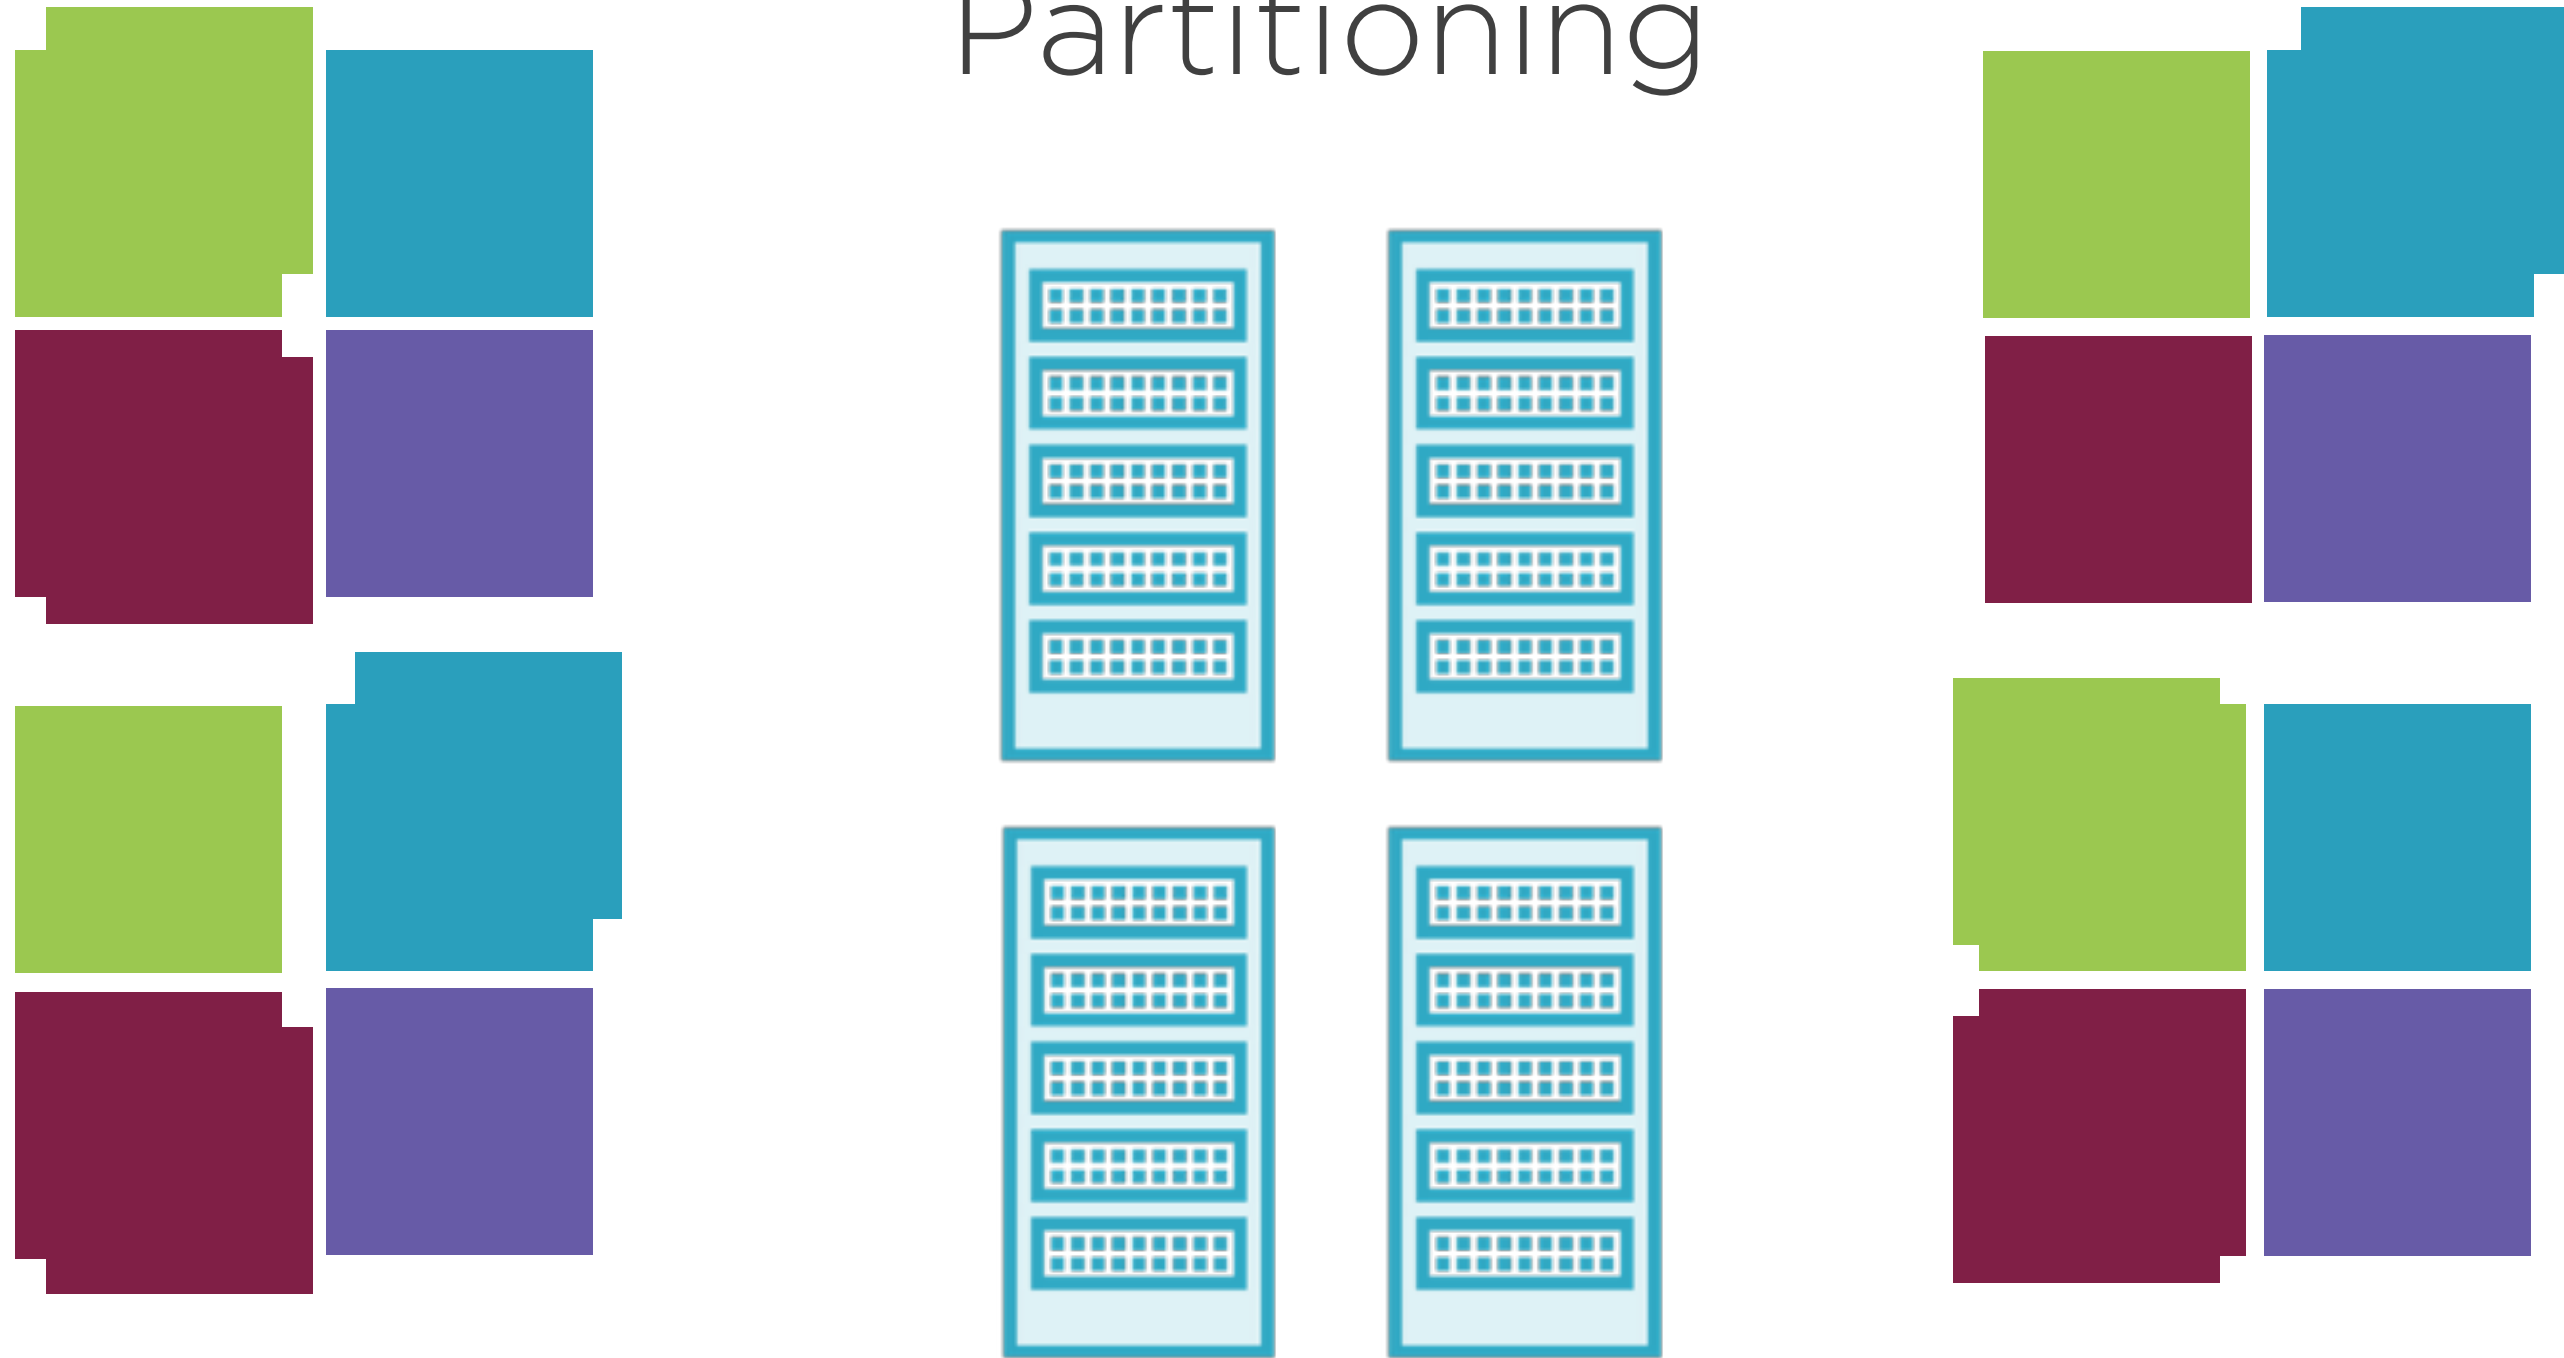
# Partitioning



**Partition this data on state information**




# Partitioning



**Records from each state will be in a different directory**

# Partitioning



State specific queries will run only on  
data in **one** directory

# Partitioning



The records in a Hive table are  
typically stored in HDFS

by default in `/user/hive/warehouse`


# Partitioning

`/user/hive/warehouse/db/orders`

`/state=CA`



`/state=NY`



`/state=WA`



`/state=NJ`



# Partitioning

**/user/hive/warehouse/db/orders**

**/state=CA**

**/file-01**

**/file-02**

**/file-03**



# Partitions vs. Indexes



## Partitions

**Physical storage of records  
affected**

**No additional data structures  
used for lookup**

**Huge datasets**

## Indexes

**No change to how the actual  
records are stored**

**Additional data structures hold a  
copy of indexed column values**

**Relatively smaller datasets**

# Demo

**Create a managed partitioned table**

**Load data into a partitioned table**

**Query the partitioned table**


**See how partitioned tables are laid out  
in HDFS**

# Why Partition Tables?

---




# Why Partition Tables?



## Logical organization of data

Set up directories and sub-directories in Hive based on how data is partitioned



## Improve query performance

Only scan the data in the partition where the result might be found

# Why Partition Tables?

ID	Customer	Product	Quantity	Amount	State
c1	John	iPhone	1	599	CA
c2	Jill	Doll	2	58	WA
c3	Emily	Shoes	1	66	NJ
c4	Nina	Jeans	1	99	NJ
c5	Rick	Skates	1	199	CA
c6	Valerie	Make Up	1	88	WA
c7	Olga	Book	3	45	NY
c8	Steven	Belt	2	25	NY

## Orders

# Why Partition Tables?

state=CA

ID	Customer	Product	Quantity	Amount	State
c1	John	iPhone	1	599	CA
c5	Rick	Skates	1	199	CA

state=WA

ID	Customer	Product	Quantity	Amount	State
c2	Jill	Doll	2	58	WA
c6	Valerie	Make	1	88	WA

state=NJ

ID	Customer	Product	Quantity	Amount	State
c3	Emily	Shoes	1	66	NJ
c4	Nina	Jeans	1	99	NJ

state=NY

ID	Customer	Product	Quantity	Amount	State
c7	Olga	Book	3	45	NY
c8	Steven	Belt	2	25	NY

```
select * from orders
where state = "NJ"
and product = "Jeans"
```

# Why Partition Tables?

**state=NJ**

ID	Customer	Product	Quantity	Amount	State
c3	Emily	Shoes	1	66	NJ
c4	Nina	Jeans	1	99	NJ

```
select * from orders  
where state = "NJ"  
and product = "Jeans"
```

Only **one partition**  
needs to be scanned

# Why Partition Tables?

**state=NJ**

ID	Customer	Product	Quantity	Amount	State
c3	Emily	Shoes	1	66	NJ
c4	Nina	Jeans	1	99	NJ

```
select * from orders  
where state = "NJ"  
and product = "Jeans"
```

**Potentially huge  
performance gains**

# Why Partition Tables?

state=CA

ID	Customer	Product	Quantity	Amount	State
c1	John	iPhone	1	599	CA
c5	Rick	Skates	1	199	CA

state=WA

ID	Customer	Product	Quantity	Amount	State
c2	Jill	Doll	2	58	WA
c6	Valerie	Make	1	88	WA

state=NJ

ID	Customer	Product	Quantity	Amount	State
c3	Emily	Shoes	1	66	NJ
c4	Nina	Jeans	1	99	NJ

state=NY

ID	Customer	Product	Quantity	Amount	State
c7	Olga	Book	3	45	NY
c8	Steven	Belt	2	25	NY

```
select * from orders
where quantity > 1
and product = "Jeans"
```

# Why Partition Tables?

state=CA

ID	Customer	Product	Quantity	Amount	State
c1	John	iPhone	1	599	CA
c5	Rick	Skates	1	199	CA

state=WA

ID	Customer	Product	Quantity	Amount	State
c2	Jill	Doll	2	58	WA
c6	Valerie	Make	1	88	WA

state=NJ

ID	Customer	Product	Quantity	Amount	State
c3	Emily	Shoes	1	66	NJ
c4	Nina	Jeans	1	99	NJ

state=NY

ID	Customer	Product	Quantity	Amount	State
c7	Olga	Book	3	45	NY
c8	Steven	Belt	2	25	NY

```
select * from orders  
where quantity > 1  
and product = "Jeans"
```

**All** partitions need to  
be scanned

# Why Partition Tables?

state=CA

ID	Customer	Product	Quantity	Amount	State
c1	John	iPhone	1	599	CA
c5	Rick	Skates	1	199	CA

state=WA

ID	Customer	Product	Quantity	Amount	State
c2	Jill	Doll	2	58	WA
c6	Valerie	Make	1	88	WA

state=NJ

ID	Customer	Product	Quantity	Amount	State
c3	Emily	Shoes	1	66	NJ
c4	Nina	Jeans	1	99	NJ

state=NY

ID	Customer	Product	Quantity	Amount	State
c7	Olga	Book	3	45	NY
c8	Steven	Belt	2	25	NY

**Partition should be based on  
the **most common queries** run**



Use **partitions** intelligently for the  
most **common** queries

# Demo

**Load data into a managed partitioned table from files**

# Demo


## **Working with an external partitioned table**

- add partitions using the alter command
- change partitions once created
- drop partitions


# Partitioning Trade-offs

---

# Partitioning Trade-offs



**How many partitions  
should you have?**




# Large Number of Partitions

Partition on customer id on the orders table?

Millions of partitions, an HDFS directory for **each** partition

**Huge overhead** for the NameNode in Hadoop

May **optimize** some queries but be **detrimental** for others



# Small Number of Partitions

Partition on product quantity on the orders table?

Very few partitions, few logical groupings in our data

**No real optimizations** on queries run


**No reduction** in the data scanned by queries

# Static and Dynamic Partitioning

---




# Static Partitioning



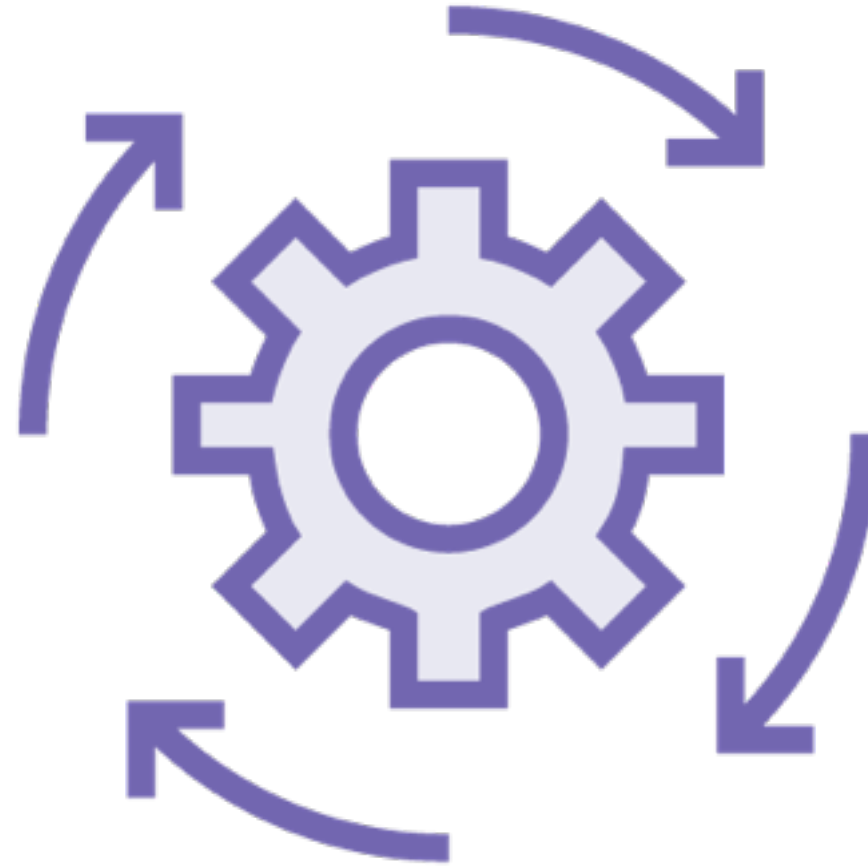
Partitions we've seen so far have been  
**static** partitions

# Static Partitioning



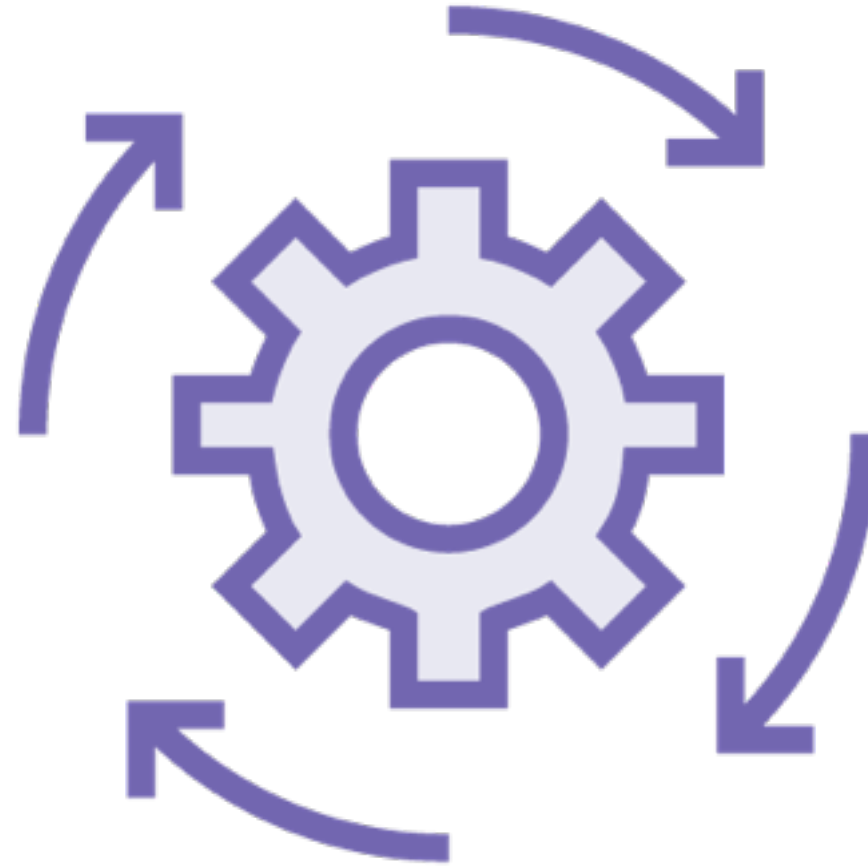
**Managing partitions manually is not a  
*scalable* process**

# Dynamic Partitioning




Dynamic partitioning allows Hive to **automatically** create partitions on data

# Dynamic Partitioning



**Partitions are based on the values of the  
partition keys**

# Dynamic Partitioning



Properties to be set in **hive-site.xml**  
or on a **per-session** basis

```
set hive.exec.dynamic.partition=true
```

---

## Enable Dynamic Partitioning

**The default for Hive is disabled, need to explicitly enable dynamic partitioning**

```
set hive.exec.dynamic.partition.mode=nonstrict
```

---

## All Partitions Can Be Dynamic

**Default is the strict mode where there should be at least one static partition**

**Non-strict allows all partitions to be dynamic**

```
set hive.exec.max.dynamic.partitions=1000
```

---

## Dynamic Partitioning Hogs Resources

**Total number of dynamic partitions created**

**If exceeded, an exception is raised at the end of the job**



```
set hive.exec.max.dynamic.partitions.pernode=3
```

---

## Dynamic Partitioning Hogs Resources

**The number of dynamic partitions created by each mapper or reducer**

**If exceeded, a fatal error will be thrown and the job will be killed**

```
set hive.exec.max.created.files=150000
```

---

## Dynamic Partitioning Hogs Resources


**The maximum number of files created by all mappers and reducers**

**If exceeded, a fatal error is thrown**


# Demo

**Set the properties to enable dynamic partitioning in Hive**

**Create and load a table with dynamic partitioning**



# Static and Dynamic Partitioning



## Static

**Partition column values are known before loading data**

**User needs to manually load data into partitions**

**Used when partitions are known and the directory structure exists in HDFS**

**Fine grained control over exact partition values**

## Dynamic

**Partition column values known only when data is loaded**

**Hive automatically loads data into partitions**

**Used when loading from an existing table that is not partitioned**

**Less control over the exact partition values**

# Demo

**Create and load a table partitioned on  
more than one column**

# Summary

**Organized data and optimized queries with Hive partitions**

**Worked with managed and external partitioned tables**

**Implemented static and dynamic partitioning and understood when each should be used**