

# Bucketing Columns for Faster Joins

---

# Overview

**Understand how Hive data can be split using bucketing**

**Know the differences between partitioning and bucketing**

**Understand the advantages of using bucketed tables**

**Learn how to implement buckets in Hive**

**Learn how to sample data from Hive tables**

# Fast Lookup of Records

---

# Fast Lookup




Requires us to know **where** exactly  
records are stored

# Hash Tables




**Data structures which allow you to lookup  
specific values very quickly**

# Hash Tables



**Have a fixed number of categories or buckets**

# Hash Tables



**Accessing a bucket is fast**


# Hash Tables




**A hash function determines which bucket  
each value belongs to**




# Hash Tables







# Hash Tables






# Hash Tables





# Hash Tables



# Hash Tables




# Hash Tables



**For any new value we know immediately  
which bucket it belongs to**


# Hash Tables



**For any new value we know immediately  
which bucket it belongs to**



# Hash Tables




Each value is **hashed** so it falls in one of these buckets





# Hash Tables



A value can only belong to **one bucket** and  
always belongs to the **same bucket**

# Hashing ~ Bucketing



**Bucketing is the **logical equivalent** of hashing in Hive**

# Hashing ~ Bucketing



**Buckets are **files** on the storage system which store those records which **map** to it**

# Hashing ~ Bucketing



**Partitions are directories, buckets are files  
under those directories**

# Bucketing Tables in Hive

---


# Bucketing



**Split data into smaller subsets**

**Separate records into manageable parts based on a column value**

# Bucketing



# Bucketing



**Records are split across multiple machines in the cluster**



# Bucketing




Each machine has directories for  
**databases, tables and partitions**

# Bucketing



And **files** which store the actual data in each directory

# Bucketing



Knowing **which file** holds the record can make lookup and join operations very fast

# Bucketing



**A hash function determines which bucket  
each value belongs to**

# Bucketing



A common hash function for integer column values is the **modulo (%)** operator

# Bucketing



**Consider an e-commerce site which sells  
millions of products**

# Bucketing


ID	Name	Cost	Category
1	iPhone	599	Mobiles
2	Doll	35	Toys
3	Shoes	33	Footwear
4	Jeans	69	Fashion
5	Skates	123	Sports
6	Make Up	99	Fashion
7	Book	24	Books
8	Belt	20	Fashion

## Products

# Bucketing

ID	Name	Cost	Category
1	iPhone	599	Mobiles
2	Doll	35	Toys
3	Shoes	33	Footwear
4	Jeans	69	Fashion
5	Skates	123	Sports
6	Make Up	99	Fashion
7	Book	24	Books
8	Belt	20	Fashion

0



1



2



bucketed on the **ID** column

Number of buckets = **3**




# Bucketing

ID	Name	Cost	Category
1	iPhone	599	Mobiles
2	Doll	35	Toys
3	Shoes	33	Footwear
4	Jeans	69	Fashion
5	Skates	123	Sports
6	Make Up	99	Fashion
7	Book	24	Books
8	Belt	20	Fashion

0



1



2



$$1 \% 3 = 1$$


# Bucketing

ID	Name	Cost	Category
1	iPhone	599	Mobiles
2	Doll	35	Toys
3	Shoes	33	Footwear
4	Jeans	69	Fashion
5	Skates	123	Sports
6	Make Up	99	Fashion
7	Book	24	Books
8	Belt	20	Fashion


0



1



2



$$2 \% 3 = 2$$

# Bucketing

ID	Name	Cost	Category
1	iPhone	599	Mobiles
2	Doll	35	Toys
3	Shoes	33	Footwear
4	Jeans	69	Fashion
5	Skates	123	Sports
6	Make Up	99	Fashion
7	Book	24	Books
8	Belt	20	Fashion

0

3

1

1

2

2

$$3 \% 3 = 0$$

# Bucketing

ID	Name	Cost	Category
1	iPhone	599	Mobiles
2	Doll	35	Toys
3	Shoes	33	Footwear
4	Jeans	69	Fashion
5	Skates	123	Sports
6	Make Up	99	Fashion
7	Book	24	Books
8	Belt	20	Fashion

0

3

1

1

4

2

2

$$4 \% 3 = 1$$

# Bucketing

ID	Name	Cost	Category
1	iPhone	599	Mobiles
2	Doll	35	Toys
3	Shoes	33	Footwear
4	Jeans	69	Fashion
5	Skates	123	Sports
6	Make Up	99	Fashion
7	Book	24	Books
8	Belt	20	Fashion

0

3

1

1

4

2

2

5

$$5 \% 3 = 2$$

# Bucketing

ID	Name	Cost	Category
1	iPhone	599	Mobiles
2	Doll	35	Toys
3	Shoes	33	Footwear
4	Jeans	69	Fashion
5	Skates	123	Sports
6	Make Up	99	Fashion
7	Book	24	Books
8	Belt	20	Fashion

0

3

6

1

1

4

2

2

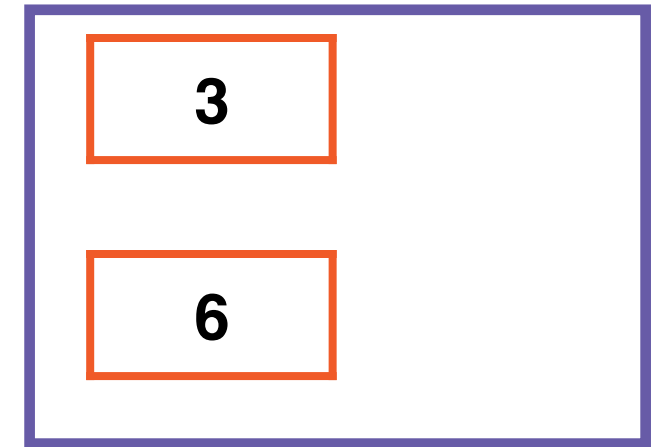
5

$$6 \% 3 = 0$$

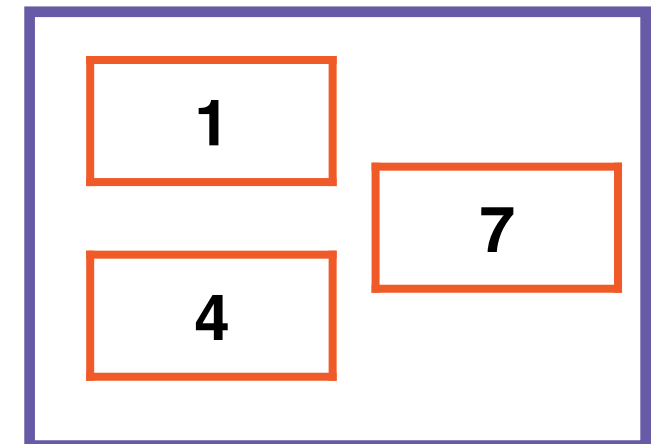
# Bucketing

ID	Name	Cost	Category
1	iPhone	599	Mobiles
2	Doll	35	Toys
3	Shoes	33	Footwear
4	Jeans	69	Fashion
5	Skates	123	Sports
6	Make Up	99	Fashion
7	Book	24	Books
8	Belt	20	Fashion


0



1



2




$$7 \% 3 = 1$$


# Bucketing

ID	Name	Cost	Category
1	iPhone	599	Mobiles
2	Doll	35	Toys
3	Shoes	33	Footwear
4	Jeans	69	Fashion
5	Skates	123	Sports
6	Make Up	99	Fashion
7	Book	24	Books
8	Belt	20	Fashion


0



1



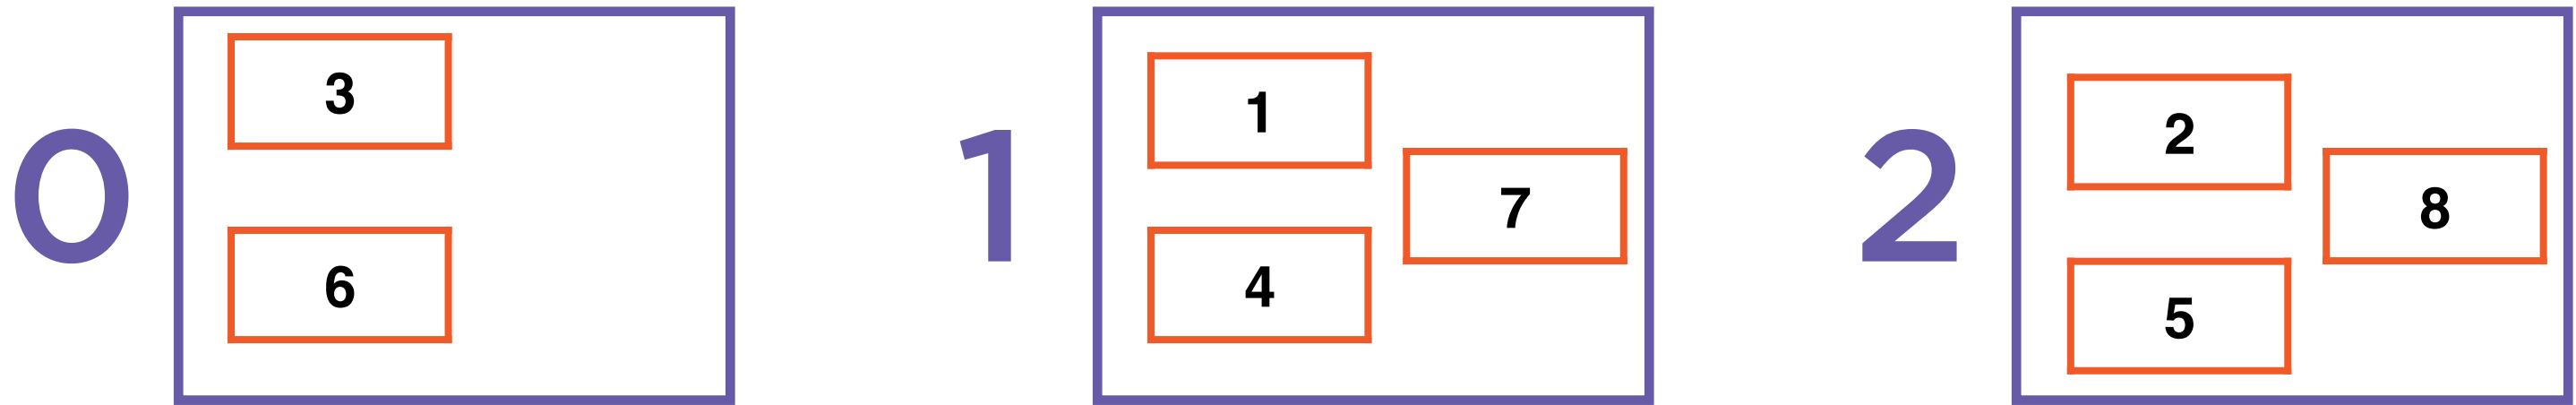
2



$$8 \% 3 = 2$$




# Bucketing



**Each of these buckets are files on HDFS**

# Bucketing



**Knowing which file each record is stored in  
makes certain operations very fast**

# Demo

**Bucketing a non-partitioned table**

**Loading data into and querying a bucketed table**

**Layout of a bucketed table**

# Bucketing vs. Partitioning

---

# Partitioning

ID		Name	Cost	Category
1		iPhone	599	Mobiles
2		Doll	35	Toys
3		Shoes	33	Footwear
4		Jeans	69	Fashion
5		Skates	123	Sports
6		Make Up	99	Fashion
7		Book	24	Books

**Partitioning on a column such as product id  
can result in **millions** of directories**

# Partitioning

ID	Name	Cost	Category
1	iPhone	599	Mobiles
2	Doll	35	Toys
3	Shoes	33	Footwear
4	Jeans	69	Fashion
5	Skates	123	Sports
6	Make Up	99	Fashion
7	Book	24	Books

Hive **restricts** the maximum number of partitions allowed to **avoid overwhelming** the NameNode

# Partitioning

ID	Name	Cost	Category
1	iPhone	599	Mobiles
2	Doll	35	Toys
3	Shoes	33	Footwear
4	Jeans	69	Fashion
5	Skates	123	Sports
6	Make Up	99	Fashion
7	Book	24	Books

**Query optimization on this column  
requires bucketing**

# Partitioning

ID	Name	Cost	Category
1	iPhone	599	Mobiles
2	Doll	35	Toys
3	Shoes	33	Footwear
4	Jeans	69	Fashion
5	Skates	123	Sports
6	Make Up	99	Fashion
7	Book	24	Books

**Fixed number** of buckets and an easy way  
to know **which** bucket holds a record





# Bucketing vs. Partitioning



## Bucketing

**Fixed** number of buckets

Based on a **hash** value of a column

Each bucket stored as **files** under a directory

Buckets are almost the **same size**

Optimizes queries such as **lookup, joins, sampling**

## Partitioning

**Unknown** number of partitions

Based on **actual** column values

Each partition stored as a **directory**

Partitions can **vary** a lot in size

Optimizes queries which **retrieve or scan** data in the logical group

# Partitioned and Bucketed Tables



**Hive tables can be both partitioned and bucketed**

# Partitioned and Bucketed Tables



**Each partitioned directory holds a file for each bucket**

# Demo

**Bucketing a partitioned table**

**Loading data into and querying a  
partitioned, bucketed table**

**Layout of a partitioned, bucketed table**

# Advantages of Bucketing

---



# Advantages of Bucketing

**Faster query responses**

**Faster join operations**

- Map side joins
- Sorted records in buckets

**More efficient sampling to test and debug operations**



# Advantages of Bucketing

**Faster query responses**

Faster join operations

- Map side joins
- Sorted records in buckets

More efficient sampling to test and debug operations

# Faster Query Responses



Given a record, it is possible to figure out  
**where exactly** the record is stored





# Advantages of Bucketing

Faster query responses

**Faster join operations**

- Map side joins
- Sorted records in buckets

More efficient sampling to test and debug operations

# Faster Join Operations



**Joins combine values from 2 or more tables  
on the **same column values****

# Faster Join Operations

## Orders

ID	Product ID	Quantity	Amount
o1	4	1	599
o2	7	1	35
o3	8	1	33
o4	5	2	69
o5	1	1	123
o6	6	1	99
o7	2	2	24
o8	3	2	20

## Products

ID	Name	Cost
7	iPhone	599
8	Doll	35
3	Shoes	33
1	Jeans	69
6	Skates	123
5	Make Up	99
4	Book	24
2	Belt	20

**Join Orders and Products to get the names of the products users have bought**

# Faster Join Operations

## Orders

ID	Product ID	Quantity	Amount
o1	4	1	599
o2	7	1	35
o3	8	1	33
o4	5	2	69
o5	1	1	123
o6	6	1	99
o7	2	2	24
o8	3	2	20

## Products

ID	Name	Cost
7	iPhone	599
8	Doll	35
3	Shoes	33
1	Jeans	69
6	Skates	123
5	Make Up	99
4	Book	24
2	Belt	20

**Product ID is the join column**

# Faster Join Operations

## Orders

ID	Product ID	Quantity	Amount
o1	<b>4</b>	1	599
o2	7	1	35
o3	8	1	33
o4	5	2	69
o5	1	1	123
o6	6	1	99
o7	2	2	24
o8	3	2	20

## Products

ID	Name	Cost
7	iPhone	599
8	Doll	35
3	Shoes	33
1	Jeans	69
6	Skates	123
5	Make Up	99
<b>4</b>	<b>Book</b>	<b>24</b>
2	Belt	20

**Need to scan the entire dataset  
to find the corresponding row**

# Faster Join Operations

## Orders

ID	Product ID	Quantity	Amount
o1	4	1	599
o2	7	1	35
o3	8	1	33
o4	5	2	69
o5	1	1	123
o6	6	1	99
o7	2	2	24
o8	3	2	20

## Products

ID	Name	Cost
6	Skates	123
3	Shoes	33
ID	Name	Cost
7	iPhone	599
1	Jeans	69
4	Book	24
ID	Name	Cost
2	Belt	20
8	Doll	35
5	Make Up	99

**Bucket the Products  
table on the ID column**

# Faster Join Operations

## Orders

ID	Product ID	Quantity	Amount
o1	4	1	599
o2	7	1	35
o3	8	1	33
o4	5	2	69
o5	1	1	123
o6	6	1	99
o7	2	2	24
o8	3	2	20

## Products

ID	Name	Cost
6	Skates	123
3	Shoes	33
ID	Name	Cost
7	iPhone	599
1	Jeans	69
<b>4</b>	<b>Book</b>	<b>24</b>
ID	Name	Cost
2	Belt	20
8	Doll	35
5	Make Up	99

**Scan a much smaller  
dataset to access each row**

# Faster Join Operations

## Orders

ID	Product ID	Quantity	Amount
o1	4	1	599
o2	7	1	35
o3	8	1	33
o4	5	2	69
o5	1	1	123
o6	6	1	99
o7	2	2	24
o8	3	2	20

## Products

ID	Name	Cost
6	Skates	123
3	Shoes	33

ID	Name	Cost
7	iPhone	599
1	Jeans	69
4	Book	24

ID	Name	Cost
2	Belt	20
8	Doll	35
5	Make Up	99

**Faster joins**





# Advantages of Bucketing

Faster query responses

**Faster join operations**

- Map side joins
- Sorted records in buckets

More efficient sampling to test and debug operations


# Joins as Map-only Operations



## MapReduce operations have 2 phases of processing



# Joins as Map-only Operations



Certain queries can be  
structured to have **no**  
**reduce** phase

# Faster Map-only Joins



**Map-only joins load one or more  
of the join tables in memory**

# Faster Map-only Joins

## Orders

ID	Product ID	Quantity	Amount
o1	4	1	599
o2	7	1	35
o3	8	1	33
o4	5	2	69
o5	1	1	123
o6	6	1	99
o7	2	2	24
o8	3	2	20

## Products

ID	Name	Cost
6	Skates	123
3	Shoes	33
ID	Name	Cost
7	iPhone	599
1	Jeans	69
4	Book	24
ID	Name	Cost
2	Belt	20
8	Doll	35
5	Make Up	99

**Mapper knows the exact bucket  
where the row is present**


# Faster Map-only Joins

## Orders

ID	Product ID	Quantity	Amount
o1	4	1	599
o2	7	1	35
o3	8	1	33
o4	5	2	69
o5	1	1	123
o6	6	1	99
o7	2	2	24
o8	3	2	20

## Products

ID	Name	Cost
6	Skates	123
3	Shoes	33
ID	Name	Cost
7	iPhone	599
1	Jeans	69
4	Book	24
ID	Name	Cost
2	Belt	20
8	Doll	35
5	Make Up	99



**Load only one bucket in  
memory**



# Advantages of Bucketing

Faster query responses

**Faster join operations**

- Map side joins
- Sorted records in buckets

More efficient sampling to test and debug operations

# Sorted Records for Faster Joins



**Records in each bucket can  
be sorted**



# Sorted Records for Faster Joins

**Orders**

ID	Product ID	Quantity	Amount
o1	4	1	599
o2	7	1	35
o3	8	1	33
o4	5	2	69
o5	1	1	123
o6	6	1	99
o7	2	2	24
o8	3	2	20

**Products**

ID	Name	Cost
6	Skates	123
3	Shoes	33
ID	Name	Cost
7	iPhone	599
1	Jeans	69
4	Book	24
ID	Name	Cost
2	Belt	20
8	Doll	35
5	Make Up	99

**Bucket both tables on the product ID column**

# Sorted Records for Faster Joins

**Orders**

ID	Product ID	Quantity	Amount
o6	6	1	99
o8	3	2	20
ID	Product ID	Quantity	Amount
o5	1	1	123
o2	7	1	35
o1	4	1	599
ID	Product ID	Quantity	Amount
o3	8	1	33
o7	2	2	24
o4	5	2	69

**Products**

ID	Name	Cost
6	Skates	123
3	Shoes	33
ID	Name	Cost
7	iPhone	599
1	Jeans	69
4	Book	24
ID	Name	Cost
2	Belt	20
8	Doll	35
5	Make Up	99

**Sort** records in each bucket on product ID

# Sorted Records for Faster Joins

## Orders

ID	Product ID	Quantity	Amount
o8	3	2	20
o6	6	1	99
ID	Product ID	Quantity	Amount
o5	1	1	123
o1	4	1	599
o2	7	1	35
ID	Product ID	Quantity	Amount
o7	2	2	24
o4	5	2	69
o3	8	1	33

## Products

ID	Name	Cost
3	Shoes	33
6	Skates	123
ID	Name	Cost
1	Jeans	69
4	Book	24
7	iPhone	599
ID	Name	Cost
2	Belt	20
5	Make Up	99
8	Doll	35

Join operations work like **merge-sort**

# Sorted Records for Faster Joins

## Orders

ID	Product ID	Quantity	Amount
o8	3	2	20
o6	6	1	99
ID	Product ID	Quantity	Amount
o5	1	1	123
o1	4	1	599
o2	7	1	35
ID	Product ID	Quantity	Amount
o7	2	2	24
o4	5	2	69
o3	8	1	33

## Products

ID	Name	Cost
3	Shoes	33
6	Skates	123
ID	Name	Cost
1	Jeans	69
4	Book	24
7	iPhone	599
ID	Name	Cost
2	Belt	20
5	Make Up	99
8	Doll	35

Faster joins

Demo

**Bucketed tables with sorted records**



# Advantages of Bucketing

Faster query responses

Faster join operations

- Map side joins
- Sorted records in buckets

**More efficient sampling to test and debug operations**

# Sampling Hive Table Data

---

# Sampling



Get a **smaller, representative subset** of data  
from the table



# Sampling

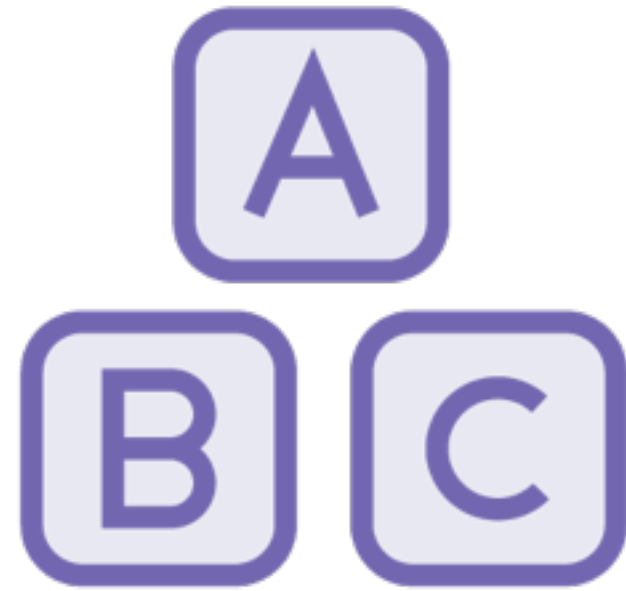


**Used for testing, aggregations, debugging**


# Sampling



Limit



Block  
sampling



Row count  
sampling



Bucket  
sampling



# Limit

Executes the query on the **entire** table

Returns only a **limited set of results** to the user

The first N rows are retrieved

**Not a random,  
representative subset**

```
select * from products limit 4;
```

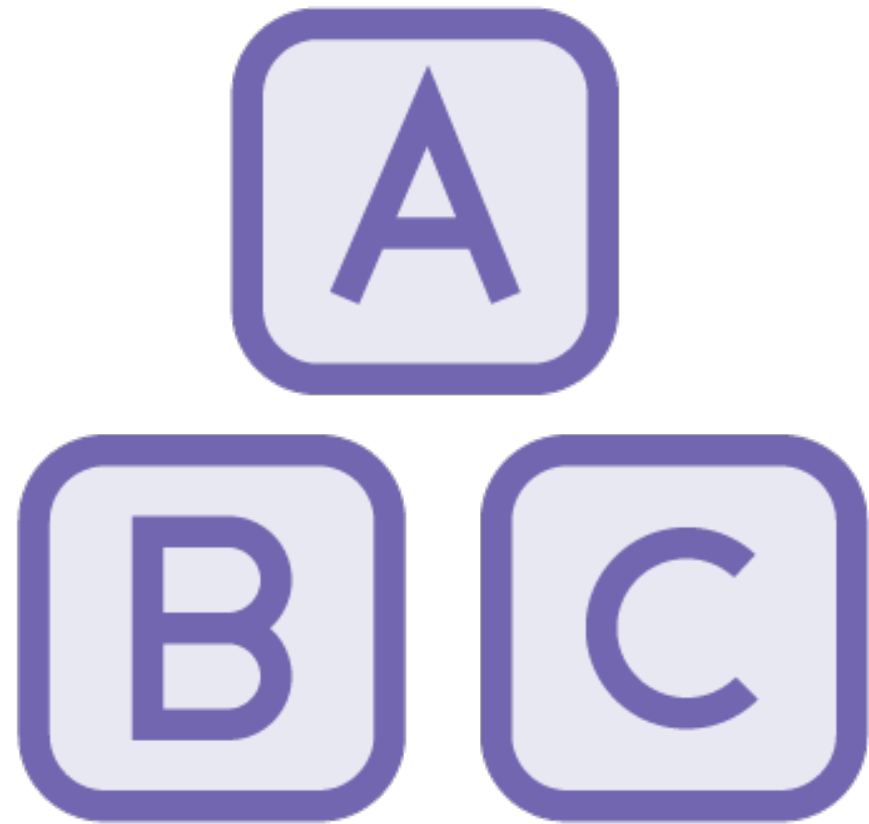
---

## Using the Limit Keyword

**Run a select on the entire table but only return 4 records to the user**

**Can be run on all tables, bucketed or non-bucketed**

# Block Sampling



Samples the **entire** table

Returns only a **few rows** to the user

A percentage of the dataset retrieved

**Inefficient way to sample  
data**

```
select * from products_no_buckets tablesample(10 percent);
```

---

## Block Sampling, Non-bucketed Tables

**Sample the entire table and return roughly 10% of the rows**


```
select * from products tablesample(10 percent);
```

---

## Block Sampling, Bucketed Tables

**Sample the entire table and return roughly 10% of the rows**

**Works on both non-buckets and bucketed tables**



# Row Count Sampling

Samples the **entire** table

Returns only a **few rows** to the user

A fixed number of rows retrieved

**Inefficient way to sample  
data**



```
select * from products_no_buckets tablesample(3 rows);
```

---

## Row Count Sampling, Non-bucketed Tables

**Sample the entire table and return exactly 3 rows**

```
select * from products tablesample(3 rows);
```

---

## Row Count Sampling, Bucketed Tables

**Sample the entire table and return exactly 3 rows**

**Works on both non-buckets and bucketed tables**



# Bucket Sampling

Samples **only a few buckets** in a table

Returns only a **few rows** to the user

A limited number of rows retrieved  
based on bucket sizes

**Much more efficient way  
to sample data**



# Bucket Sampling

Samples **only a few buckets** in a table

Returns only a **few rows** to the user

A limited number of rows retrieved  
based on bucket sizes

**Works only on bucketed  
tables**

Demo

**Sample data in Hive tables using bucket sampling**

# Summary

**Understood how bucketing works and the advantages of bucketing in Hive**

**Understood the differences between partitioning and bucketing**

**Learnt how to implement buckets in Hive**

**Learnt how to sample data from Hive tables**