

# Optimizing Hive Joins

---

# Overview

**Understand how joins are implemented in Hive**

**Optimize joins which involve large tables**

**Use semi-joins in place of IN/EXISTS subqueries**

**Optimize map-only joins**

# Join Operations as MapReduce Jobs


---

# Join Operations



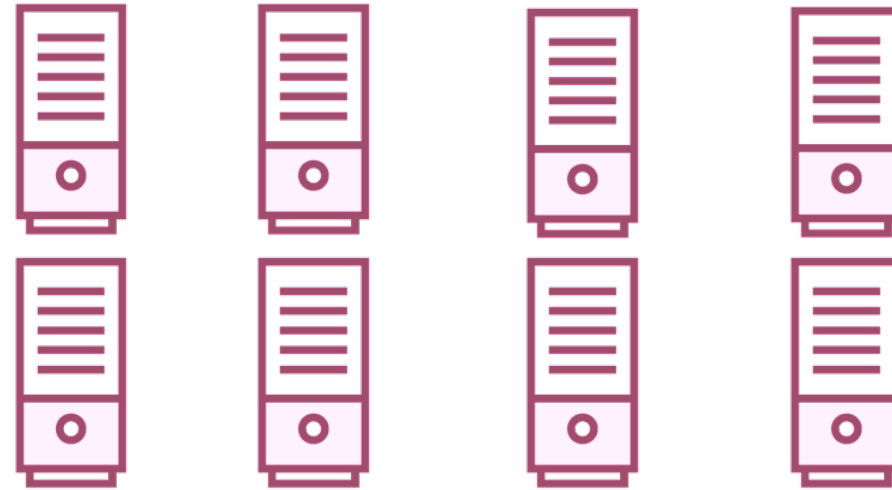
**Join operations are MapReduce  
jobs under the hood**

# MapReduce




**A programming paradigm which  
runs on a distributed system**

# MapReduce



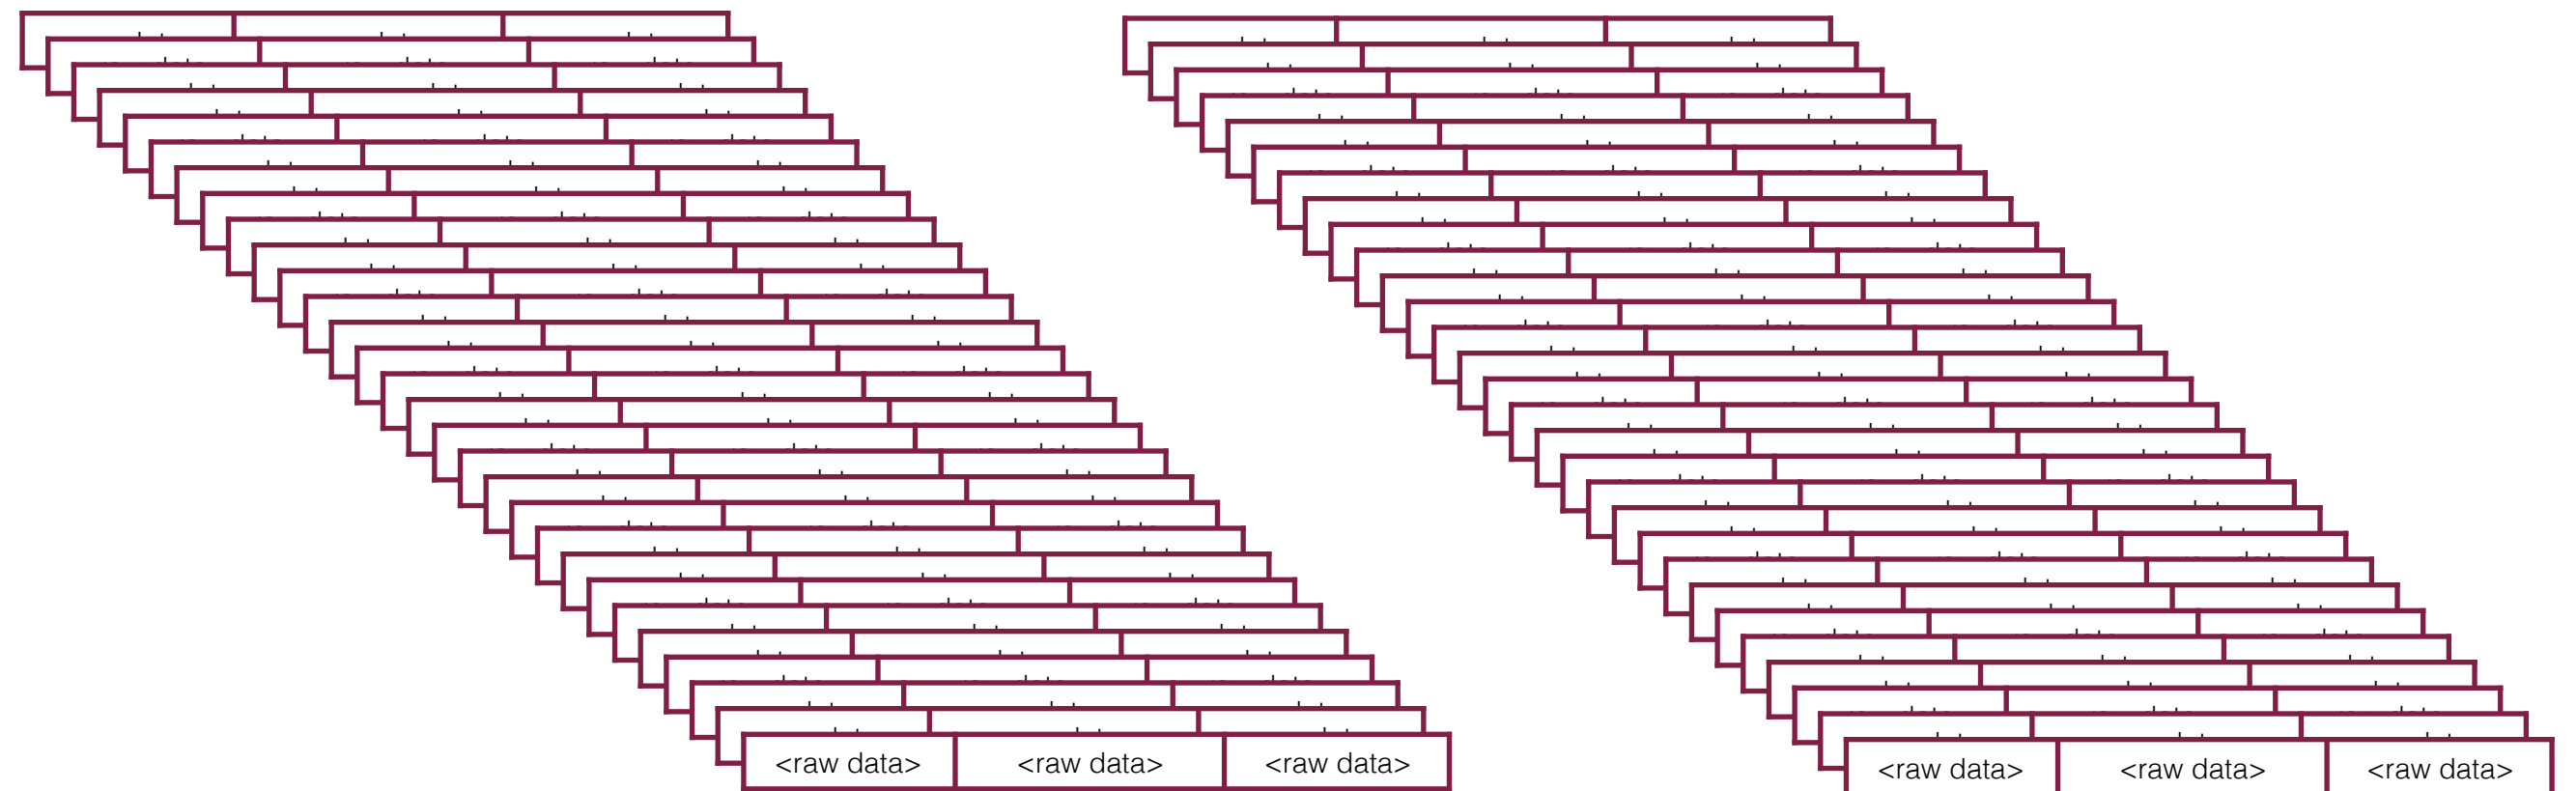
**Takes advantage of the inherent  
parallelism in data processing**

# MapReduce



**Modern systems generate millions of  
records of raw data**

# MapReduce







A task of this scale is processed in  
two stages

map

reduce



# map



# reduce



|            |            |            |
|------------|------------|------------|
| <raw data> | <raw data> | <raw data> |
| <raw data> | <raw data> | <raw data> |
| <raw data> | <raw data> | <raw data> |
| <raw data> | <raw data> | <raw data> |




# map

**An operation performed  
in parallel, on small  
portions of the dataset**

# map

One Record




Key-Value Output

# reduce

**An operation to  
combine the results of  
the map step**

# reduce



**map**

A step that can be performed in parallel

**reduce**

A step to combine the intermediate results

# Join Columns and MapReduce Jobs

---



# Join Columns



Join combines records from two or more tables on the **same column value**

# Join Columns

## Trades

| Symbol | Open | High | Low | Close | Day      |
|--------|------|------|-----|-------|----------|
| GOOG   | 820  | 840  | 818 | 829   | 1-1-2017 |

## Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

## Revenues

| Symbol | Name      | Revenue |
|--------|-----------|---------|
| GOOG   | Google    | 90B     |
| AAPL   | Apple     | 215B    |
| MSFT   | Microsoft | 85B     |

```
select * from Names join Trades
on Names.Symbol = Trades.Symbol
```

# Join Columns

## Trades

| Symbol | Open | High | Low | Close | Day      |
|--------|------|------|-----|-------|----------|
| GOOG   | 820  | 840  | 818 | 829   | 1-1-2017 |

## Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

## Revenues

| Symbol | Name      | Revenue |
|--------|-----------|---------|
| GOOG   | Google    | 90B     |
| AAPL   | Apple     | 215B    |
| MSFT   | Microsoft | 85B     |

**One** join column = **one** MapReduce job

# Join Columns

## Trades

| Symbol | Open | High | Low | Close | Day      |
|--------|------|------|-----|-------|----------|
| GOOG   | 820  | 840  | 818 | 829   | 1-1-2017 |

## Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

## Revenues

| Symbol | Name      | Revenue |
|--------|-----------|---------|
| GOOG   | Google    | 90B     |
| AAPL   | Apple     | 215B    |
| MSFT   | Microsoft | 85B     |

```
select * from Names join Trades
on (Names.Symbol = Trades.Symbol)
join Revenues on (Names.Symbol = Revenues.Symbol)
```

# Join Columns

## Trades

| Symbol | Open | High | Low | Close | Day      |
|--------|------|------|-----|-------|----------|
| GOOG   | 820  | 840  | 818 | 829   | 1-1-2017 |

## Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

## Revenues

| Symbol | Name      | Revenue |
|--------|-----------|---------|
| GOOG   | Google    | 90B     |
| AAPL   | Apple     | 215B    |
| MSFT   | Microsoft | 85B     |

**One** join column = **one** MapReduce job

# Join Columns

## Trades

| Symbol | Open | High | Low | Close | Day      |
|--------|------|------|-----|-------|----------|
| GOOG   | 820  | 840  | 818 | 829   | 1-1-2017 |

## Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

## Revenues

| Symbol | Name      | Revenue |
|--------|-----------|---------|
| GOOG   | Google    | 90B     |
| AAPL   | Apple     | 215B    |
| MSFT   | Microsoft | 85B     |

```
select * from Names join Trades
on (Names.Symbol = Trades.Symbol)
join Revenues on (Names.Name = Revenues.Name)
```

# Join Columns

## Trades

| Symbol | Open | High | Low | Close | Day      |
|--------|------|------|-----|-------|----------|
| GOOG   | 820  | 840  | 818 | 829   | 1-1-2017 |

## Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

## Revenues

| Symbol | Name      | Revenue |
|--------|-----------|---------|
| GOOG   | Google    | 90B     |
| AAPL   | Apple     | 215B    |
| MSFT   | Microsoft | 85B     |

```
select * from Names join Trades
on (Names.Symbol = Trades.Symbol)
join Revenues on (Names.Name = Revenues.Name)
```

# Join Columns

## Trades

| Symbol | Open | High | Low | Close | Day      |
|--------|------|------|-----|-------|----------|
| GOOG   | 820  | 840  | 818 | 829   | 1-1-2017 |

## Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

## Revenues

| Symbol | Name      | Revenue |
|--------|-----------|---------|
| GOOG   | Google    | 90B     |
| AAPL   | Apple     | 215B    |
| MSFT   | Microsoft | 85B     |

**Two** join columns = **two** MapReduce jobs



For faster queries...

**Minimize** the number of  
MapReduce jobs run

# Demo

**Join operations on 3 tables, with  
different columns in the join clause**

# Join Operations and Table Sizes

---

# Size of Tables

**Trades 500GB**

| Symbol | Open | High | Low | Close |
|--------|------|------|-----|-------|
| GOOG   | 820  | 840  | 818 | 829   |

**Names 10MB**

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

**Revenues 100MB**

| Symbol | Name      | Revenue |
|--------|-----------|---------|
| GOOG   | Google    | 90B     |
| AAPL   | Apple     | 215B    |
| MSFT   | Microsoft | 85B     |

```
select * from Names join Trades
on (Names.Symbol = Trades.Symbol)
join Revenues on (Names.Symbol = Revenues.Symbol)
```

# Size of Tables

**Trades 500GB**

| Symbol | Open | High | Low | Close |
|--------|------|------|-----|-------|
| GOOG   | 820  | 840  | 818 | 829   |

**Names 10MB**

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

**Revenues 100MB**

| Symbol | Name      | Revenue |
|--------|-----------|---------|
| GOOG   | Google    | 90B     |
| AAPL   | Apple     | 215B    |
| MSFT   | Microsoft | 85B     |

```
select * from Names join Trades
on (Names.Symbol = Trades.Symbol)
join Revenues on (Names.Symbol = Revenues.Symbol)
```

# Size of Tables

Trades **500GB**

| Symbol | Open | High | Low | Close |
|--------|------|------|-----|-------|
| GOOG   | 820  | 840  | 818 | 829   |

Names **10MB**

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

Revenues **100MB**

| Symbol | Name      | Revenue |
|--------|-----------|---------|
| GOOG   | Google    | 90B     |
| AAPL   | Apple     | 215B    |
| MSFT   | Microsoft | 85B     |

Names  
**10MB**

Trades  
**500GB**

Revenues  
**100MB**

# Size of Tables

**Names**

**10MB**

**Trades**

**500GB**

Revenues

100MB



All tables, except for the  
last are held in **memory**

# Size of Tables

Names


10MB

Trades

500GB

Revenues

100MB



The last table is **streamed**  
from disk to the job



# Size of Tables

**Names**

**10MB**

**Trades**

**500GB**

**Revenues**

**100MB**



**Inefficient to keep large  
tables in memory**

# Size of Tables

**Names**


**10MB**

**Revenues**

**100MB**

**Trades**

**500GB**



Re-order the tables in the join  
so the **largest** table is at the end

# Size of Tables

**Trades 500GB**

| Symbol | Open | High | Low | Close |
|--------|------|------|-----|-------|
| GOOG   | 820  | 840  | 818 | 829   |

**Names 10MB**

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

**Revenues 100MB**

| Symbol | Name      | Revenue |
|--------|-----------|---------|
| GOOG   | Google    | 90B     |
| AAPL   | Apple     | 215B    |
| MSFT   | Microsoft | 85B     |

```
select * from Names join Revenues
on (Names.Symbol = Revenues.Symbol)
join Trades on (Names.Symbol = Trades.Symbol)
```

# Size of Tables

**Trades 500GB**

| Symbol | Open | High | Low | Close |
|--------|------|------|-----|-------|
| GOOG   | 820  | 840  | 818 | 829   |

**Names 10MB**

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

**Revenues 100MB**

| Symbol | Name      | Revenue |
|--------|-----------|---------|
| GOOG   | Google    | 90B     |
| AAPL   | Apple     | 215B    |
| MSFT   | Microsoft | 85B     |

```
select * from Names join Revenues
on (Names.Symbol = Revenues.Symbol)
join Trades on (Names.Symbol = Trades.Symbol)
```

For faster queries...

Specify the **largest** table  
at the very **end**

```
select /*+ streamtable(Trades) */  
Names.Symbol, Trades.High, Revenues.Revenue  
from Names join Trades  
on (Names.Symbol = Trades.Symbol)  
join Revenues on (Names.Symbol = Revenues.Symbol)
```

---

## The Streamtable Keyword

**Specify which table to stream in a join operation**

**Stream the largest table, do not hold it in memory**

# Join Optimizations with Bucketing and Partitioning

---

# Bucketing

## Orders

| ID | Product ID | Quantity | Amount |
|----|------------|----------|--------|
| o1 | 4          | 1        | 599    |
| o2 | 7          | 1        | 35     |
| o3 | 8          | 1        | 33     |
| o4 | 5          | 2        | 69     |
| o5 | 1          | 1        | 123    |
| o6 | 6          | 1        | 99     |
| o7 | 2          | 2        | 24     |
| o8 | 3          | 2        | 20     |

## Products

| ID | Name    | Cost |
|----|---------|------|
| 1  | iPhone  | 599  |
| 2  | Doll    | 35   |
| 3  | Shoes   | 33   |
| 4  | Jeans   | 69   |
| 5  | Skates  | 123  |
| 6  | Make Up | 99   |
| 7  | Book    | 24   |
| 8  | Belt    | 20   |

**Join Orders and Products to get the names of the products users have bought**



# Bucketing

## Orders

| ID | Product ID | Quantity | Amount |
|----|------------|----------|--------|
| o1 | 4          | 1        | 599    |
| o2 | 7          | 1        | 35     |
| o3 | 8          | 1        | 33     |
| o4 | 5          | 2        | 69     |
| o5 | 1          | 1        | 123    |
| o6 | 6          | 1        | 99     |
| o7 | 2          | 2        | 24     |
| o8 | 3          | 2        | 20     |

## Products

| ID | Name    | Cost |
|----|---------|------|
| 7  | iPhone  | 599  |
| 8  | Doll    | 35   |
| 3  | Shoes   | 33   |
| 1  | Jeans   | 69   |
| 6  | Skates  | 123  |
| 5  | Make Up | 99   |
| 4  | Book    | 24   |
| 2  | Belt    | 20   |

**Product ID is the join column**

# Bucketing

## Orders

| ID | Product ID | Quantity | Amount |
|----|------------|----------|--------|
| o1 | <b>4</b>   | 1        | 599    |
| o2 | 7          | 1        | 35     |
| o3 | 8          | 1        | 33     |
| o4 | 5          | 2        | 69     |
| o5 | 1          | 1        | 123    |
| o6 | 6          | 1        | 99     |
| o7 | 2          | 2        | 24     |
| o8 | 3          | 2        | 20     |

## Products

| ID       | Name        | Cost      |
|----------|-------------|-----------|
| 7        | iPhone      | 599       |
| 8        | Doll        | 35        |
| 3        | Shoes       | 33        |
| 1        | Jeans       | 69        |
| 6        | Skates      | 123       |
| 5        | Make Up     | 99        |
| <b>4</b> | <b>Book</b> | <b>24</b> |
| 2        | Belt        | 20        |

Need to scan the entire dataset  
to find the corresponding row

# Bucketing

## Orders

| ID | Product ID | Quantity | Amount |
|----|------------|----------|--------|
| o1 | 4          | 1        | 599    |
| o2 | 7          | 1        | 35     |
| o3 | 8          | 1        | 33     |
| o4 | 5          | 2        | 69     |
| o5 | 1          | 1        | 123    |
| o6 | 6          | 1        | 99     |
| o7 | 2          | 2        | 24     |
| o8 | 3          | 2        | 20     |

## Products

| ID | Name    | Cost |
|----|---------|------|
| 6  | Skates  | 123  |
| 3  | Shoes   | 33   |
| ID | Name    | Cost |
| 7  | iPhone  | 599  |
| 1  | Jeans   | 69   |
| 4  | Book    | 24   |
| ID | Name    | Cost |
| 2  | Belt    | 20   |
| 8  | Doll    | 35   |
| 5  | Make Up | 99   |

**Bucket the Products  
table on the ID column**

# Bucketing

## Orders

| ID | Product ID | Quantity | Amount |
|----|------------|----------|--------|
| o1 | 4          | 1        | 599    |
| o2 | 7          | 1        | 35     |
| o3 | 8          | 1        | 33     |
| o4 | 5          | 2        | 69     |
| o5 | 1          | 1        | 123    |
| o6 | 6          | 1        | 99     |
| o7 | 2          | 2        | 24     |
| o8 | 3          | 2        | 20     |

## Products

| ID       | Name        | Cost      |
|----------|-------------|-----------|
| 6        | Skates      | 123       |
| 3        | Shoes       | 33        |
| ID       | Name        | Cost      |
| 7        | iPhone      | 599       |
| 1        | Jeans       | 69        |
| <b>4</b> | <b>Book</b> | <b>24</b> |
| ID       | Name        | Cost      |
| 2        | Belt        | 20        |
| 8        | Doll        | 35        |
| 5        | Make Up     | 99        |

**Scan a much smaller  
dataset to access each row**

# Bucketing

## Orders

| ID | Product ID | Quantity | Amount |
|----|------------|----------|--------|
| o1 | 4          | 1        | 599    |
| o2 | 7          | 1        | 35     |
| o3 | 8          | 1        | 33     |
| o4 | 5          | 2        | 69     |
| o5 | 1          | 1        | 123    |
| o6 | 6          | 1        | 99     |
| o7 | 2          | 2        | 24     |
| o8 | 3          | 2        | 20     |

## Products

| ID | Name    | Cost |
|----|---------|------|
| 6  | Skates  | 123  |
| 3  | Shoes   | 33   |
| ID | Name    | Cost |
| 7  | iPhone  | 599  |
| 1  | Jeans   | 69   |
| 4  | Book    | 24   |
| ID | Name    | Cost |
| 2  | Belt    | 20   |
| 8  | Doll    | 35   |
| 5  | Make Up | 99   |

**Faster joins**

# Partitioning

**Join optimizations would work the same way**

**Reduce** the dataset to scan to find the **corresponding row**

For faster queries...

Use **bucketing** or **partitioning**  
on the join columns

# Left Semi-joins Instead of Subqueries

---



# Left Semi-join

Any join query that requests rows from the left row source based on the existence of rows in the right row source without including data from the right row source in the final result and without duplicating rows from the left row source is a logical left semi join.

<http://sqlity.net/>

# Left Semi-join

Any join query that requests rows from the left row source based on the existence of rows in the right row source without including data from the right row source in the final result and without duplicating rows from the left row source is a logical left semi join.

<http://sqlity.net/>

# Left Semi-join

Any join query that requests rows from the left row source based on the existence of rows in the right row source without including data from the right row source in the final result and without duplicating rows from the left row source is a logical left semi join.

<http://sqlity.net/>

# Left Semi-join

Any join query that requests rows from the left row source based on the existence of rows in the right row source without including data from the right row source in the final result and **without duplicating rows from the left row source** is a logical left semi join.

<http://sqlity.net/>

# Left Semi-join

**Names**

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

**Trades**

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |
| MSFT   |      |      |     |       |     |
|        |      |      |     |       |     |

```
select names.symbol
from names left semi join trades
on
(names.symbol = trades.symbol);
```

# Left Semi-join

**Names**

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

**Trades**

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |
| MSFT   |      |      |     |       |     |
|        |      |      |     |       |     |

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| MSFT   | Microsoft |

```
select symbol from names  
where symbol in  
(select symbol from trades);
```

---

## The IN/NOT IN Keywords

**Check for existence of a column value in the subquery**

```
select symbol from names  
where symbol in  
(select symbol from trades);
```

---

## The IN/NOT IN Keywords

**Only one column can be selected in the subquery**



```
select names.symbol  
from names  
where exists  
(select trades.symbol from trades  
where names.symbol = trades.symbol);
```

---

## The EXISTS/NOT EXISTS Keywords

**Check if even one result is returned from the subquery**

```
select names.symbol  
from names  
where exists  
(select trades.symbol from trades  
where names.symbol = trades.symbol);
```

---

## The EXISTS/NOT EXISTS Keywords

**The subquery has to reference the parent query i.e. should be correlated**

# Left Semi-join

**It is much **more efficient** to  
replace the IN/EXISTS subqueries  
with a **left semi-join****

```
select names.symbol  
from names left semi join trades  
on  
(names.symbol = trades.symbol);
```

---

## The Left Semi-join Operator

**Choose records from the left table**

```
select names.symbol  
from names left semi join trades  
on  
(names.symbol = trades.symbol);
```

---

## The Left Semi-join Operator

**Choose records from the left table**

```
select names.symbol  
from names left semi join trades  
on  
(names.symbol = trades.symbol);
```

---

## The Left Semi-join Operator

**Choose records from the left table**

**Based on the existence of rows in the right table**

```
select names.symbol  
from names left semi join revenues  
on  
(names.symbol = revenues.symbol  
and  
names.name = revenues.name);
```

---

## The Left Semi-join Operator

**Allows specifying multiple columns in the join**

**Not allowed when we use the IN keyword**

# Left Semi-join vs. In/Exists

## Left Semi-join

Only scans the right table till **a match** is found

Can specify matches on **multiple** columns

**Bucketing** or **partitioning** the right table can improve performance

## In/Exists

Needs to scan the **entire** subquery table

IN allows **only one column** to be selected in the subquery

**No advantage** of a bucketed or partitioned right table



For faster queries...

Use the **left semi-join** rather than in/exists keywords with subqueries

# Demo

**The left semi-join as opposed to the in/  
exists subqueries**


# The Anatomy of a MapReduce Program

---

# Counting Word Frequencies

## Consider a large text file

|                             |
|-----------------------------|
| Twinkle twinkle little star |
| How I wonder what you are   |
| Up above the world so high  |
| Like a diamond in the sky   |
| Twinkle twinkle little star |
| How I wonder what you are   |
| .....                       |



| Word    | Frequency |
|---------|-----------|
| above   | 14        |
| are     | 20        |
| how     | 21        |
| star    | 22        |
| twinkle | 32        |
| ...     | ..        |

# MapReduce Flow

|                             |
|-----------------------------|
| Twinkle twinkle little star |
| How I wonder what you are   |



|                            |
|----------------------------|
| Up above the world so high |
| Like a diamond in the sky  |


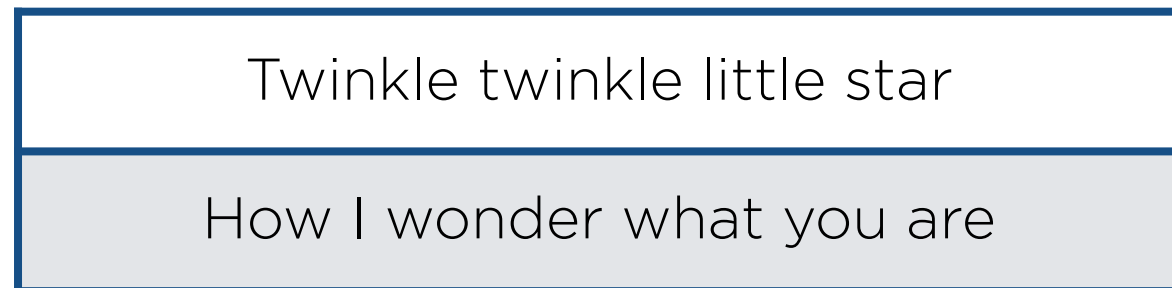

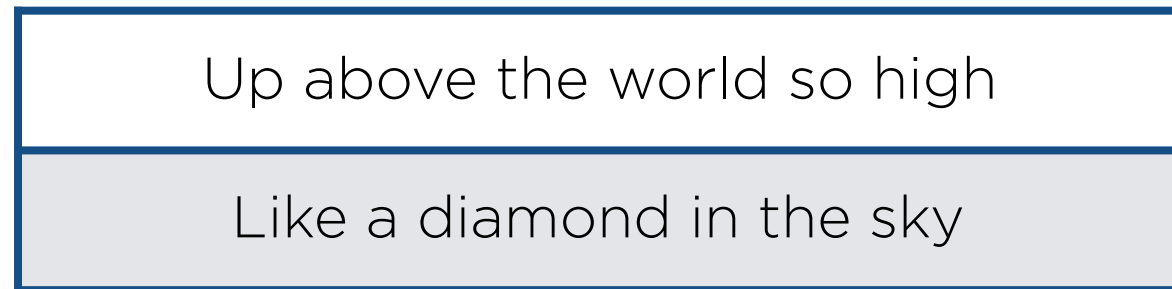

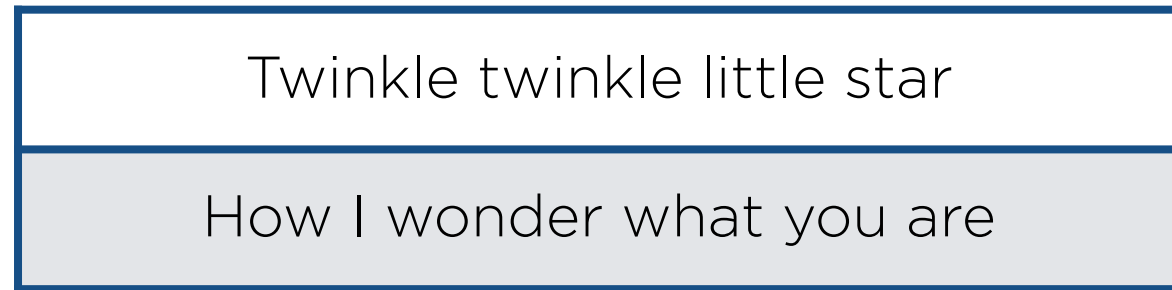


|                             |
|-----------------------------|
| Twinkle twinkle little star |
| How I wonder what you are   |



**Each partition is given to a  
different process i.e. to  
mappers**


# MapReduce Flow



**Each mapper  
works in parallel**

# Map Flow


|                             |
|-----------------------------|
| Twinkle twinkle little star |
| How I wonder what you are   |



**Within each mapper, the rows  
are processed serially**

# Map Flow

|                             |
|-----------------------------|
| Twinkle twinkle little star |
| How I wonder what you are   |



| Word | # Count |
|------|---------|
|------|---------|

{twinkle, 1}

{twinkle, 1}


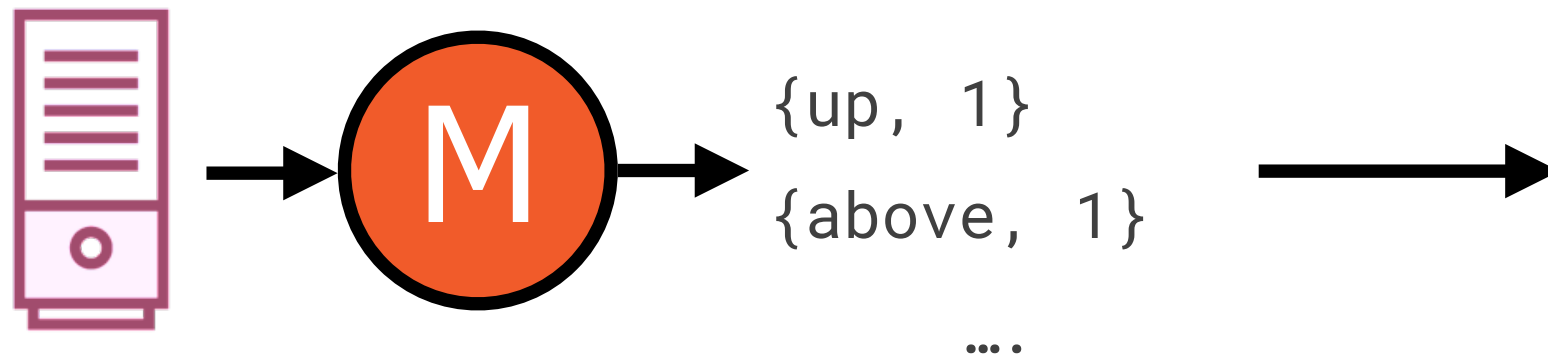

{little, 1}

{star, 1}

**Each row emits {key, value} pairs**




# Reduce Flow



**The results are  
passed on to another  
process i.e. a reducer**

# Reduce Flow



# Key Insight Behind MapReduce



Many data processing tasks can be expressed in this form

# MapReduce

## Orders

| ID | Product ID | Quantity | Amount |
|----|------------|----------|--------|
| o1 | 4          | 1        | 599    |
| o2 | 7          | 1        | 35     |
| o3 | 8          | 1        | 33     |
| o4 | 5          | 2        | 69     |
| o5 | 1          | 1        | 123    |
| o6 | 6          | 1        | 99     |
| o7 | 2          | 2        | 24     |
| o8 | 3          | 2        | 20     |

## Products

| ID | Name    | Cost |
|----|---------|------|
| 1  | iPhone  | 599  |
| 2  | Doll    | 35   |
| 3  | Shoes   | 33   |
| 4  | Jeans   | 69   |
| 5  | Skates  | 123  |
| 6  | Make Up | 99   |
| 7  | Book    | 24   |
| 8  | Belt    | 20   |

```
select * from Orders join Products
on Orders.ProductID = Products.ID
```

# MapReduce

## Orders

| ID | Product ID | Quantity | Amount |
|----|------------|----------|--------|
| o1 | 4          | 1        | 599    |
| o2 | 7          | 1        | 35     |



## Products

| ID | Name   | Cost |
|----|--------|------|
| 1  | iPhone | 599  |
| 2  | Doll   | 35   |



The mapper operates on each row of the tables

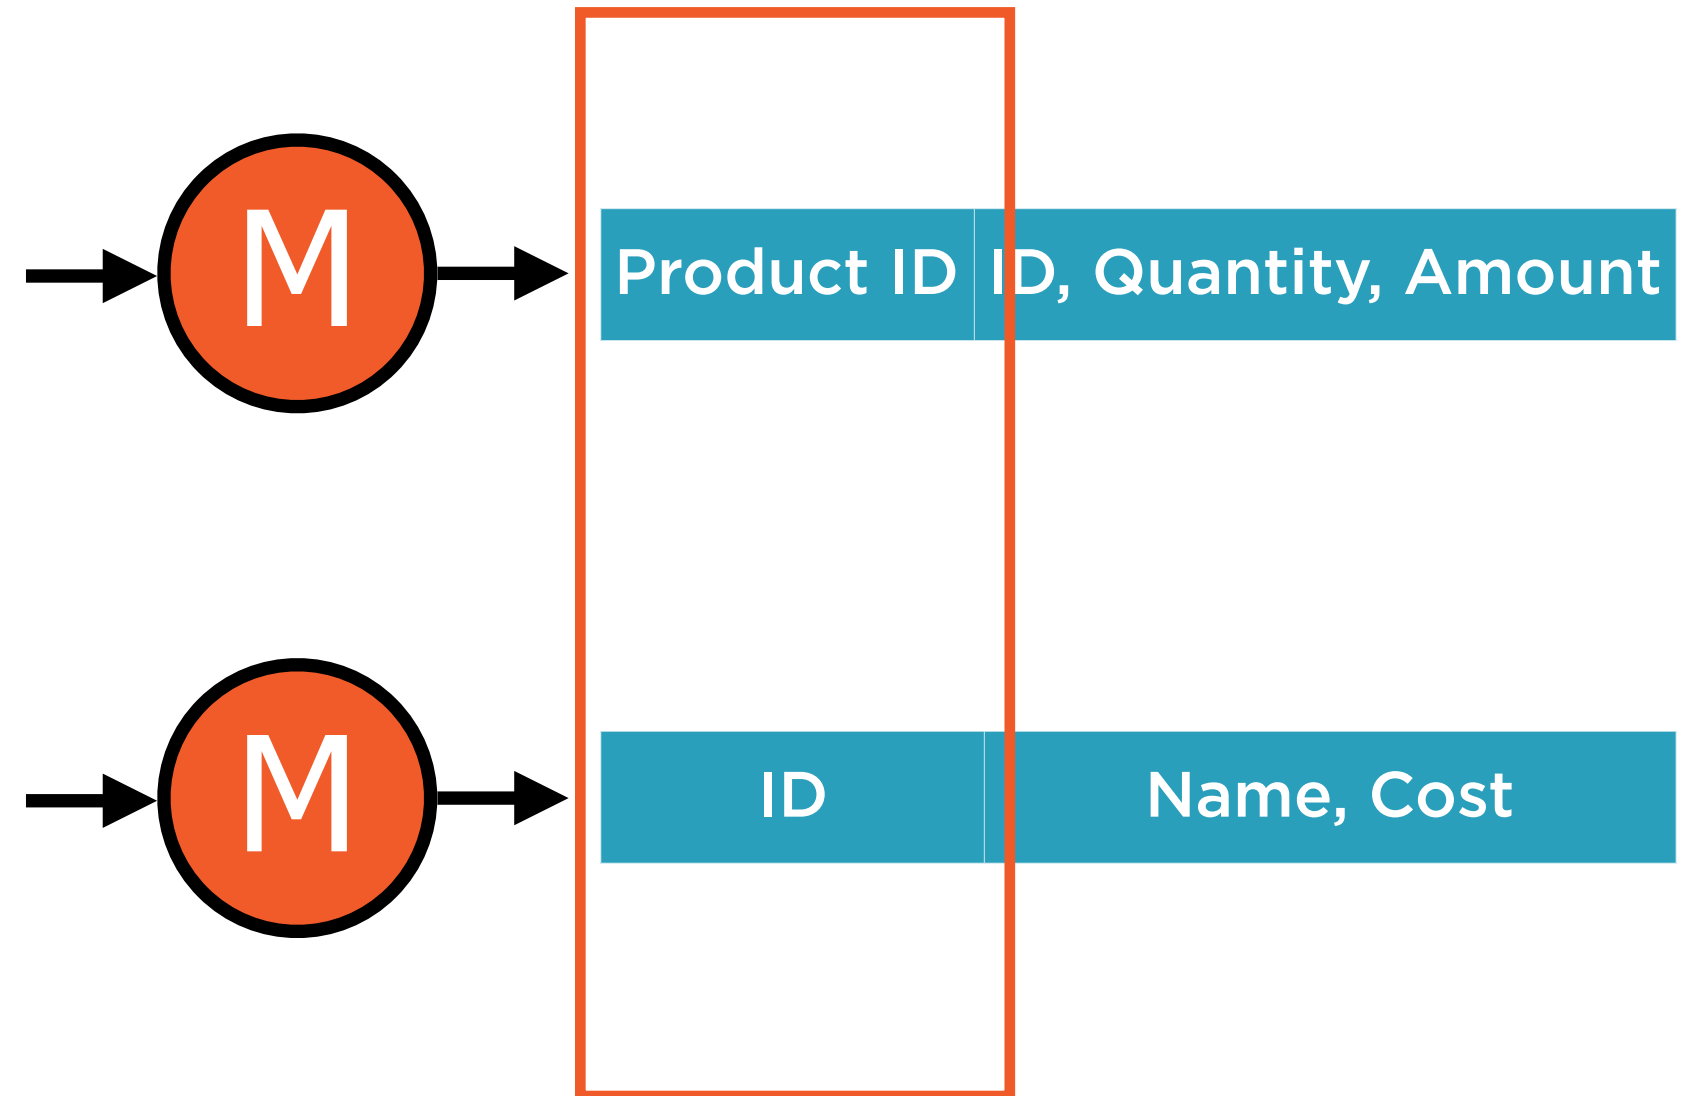
# MapReduce

## Orders

| ID | Product ID | Quantity | Amount |
|----|------------|----------|--------|
| o1 | 4          | 1        | 599    |
| o2 | 7          | 1        | 35     |

## Products

| ID | Name   | Cost |
|----|--------|------|
| 1  | iPhone | 599  |
| 2  | Doll   | 35   |



The join column is the key


# MapReduce

## Orders


| ID | Product ID | Quantity | Amount |
|----|------------|----------|--------|
| o1 | 4          | 1        | 599    |
| o2 | 7          | 1        | 35     |

## Products

| ID | Name   | Cost |
|----|--------|------|
| 1  | iPhone | 599  |
| 2  | Doll   | 35   |



Product ID | ID, Quantity, Amount

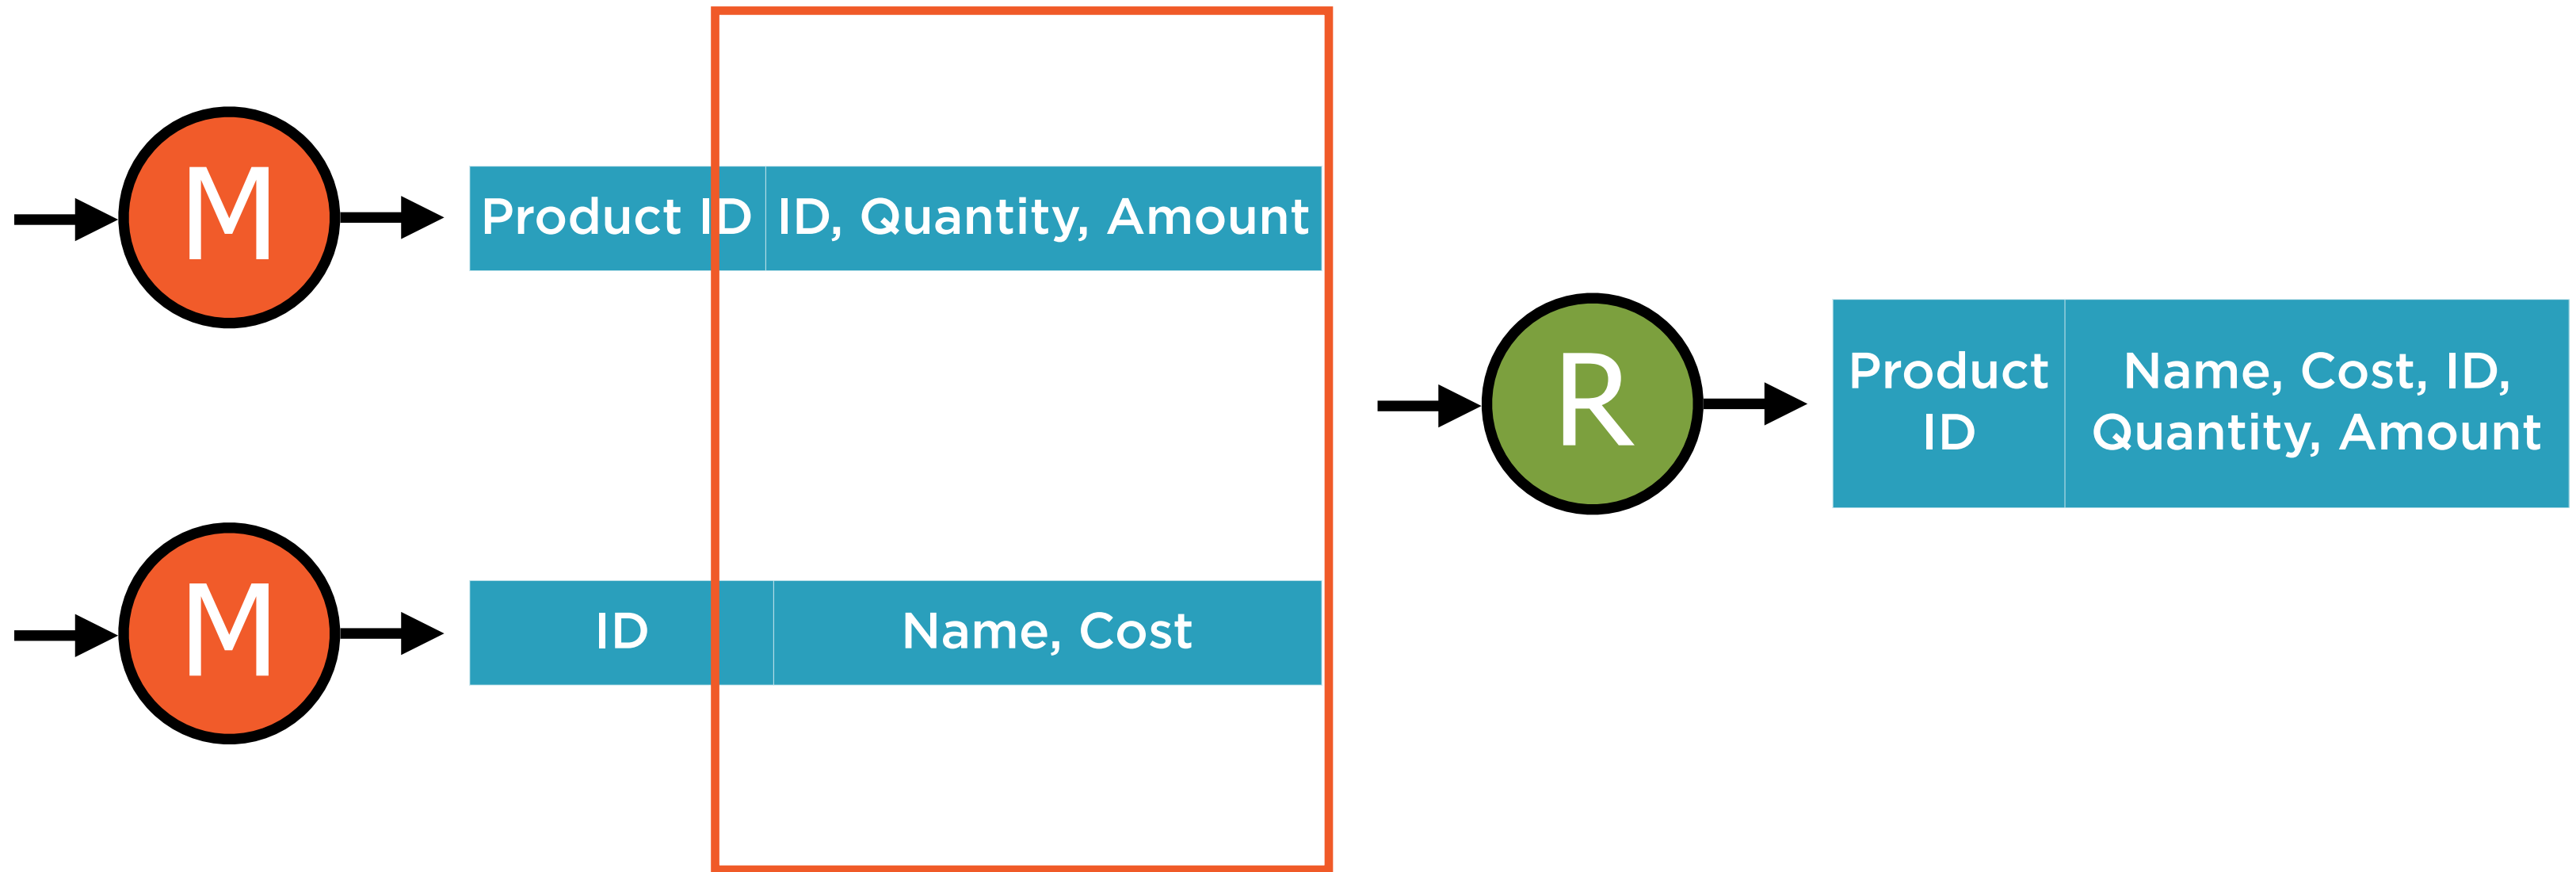


ID

Name, Cost

The remaining columns are  
values

# MapReduce



**The reducer combines all columns which have the same key**




# Joins as Map-only Operations



## MapReduce operations have 2 phases of processing




# Joins as Map-only Operations



Certain queries can be  
structured to have **no**  
**reduce** phase


# Joins as Map-only Operations



Such joins are called  
map-side joins

**More performant**

# Joins as Map-only Operations




Improves processing time

Reduces data transfer between machines in the cluster

Reduces operations such as shuffle and sort between map and reduce phases

# Joins as Map-only Operations



**We'd like joins to be  
map-side joins if possible**

# Conditions for Map-side Joins

---


# Conditions for Map-side Joins



All tables, except one, are **small** enough to be held in **memory**

Tables are **bucketed on the join columns** and  
 $\text{Table1 Buckets} = N * \text{Table2 Buckets}$

# Conditions for Map-side Joins



All tables, except one, are **small** enough to be held in **memory**

Tables are bucketed on the join columns and  
 $\text{Table1 Buckets} = N * \text{Table2 Buckets}$



# All Tables but One Are Small

**Names**

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

**Trades**

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |

Left table is the  
smaller table

# Inner Join

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

Trades

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |

Only rows which have a match  
in **both** the left and right table

# Inner Join

**Names**

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

**Trades**

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |

```
select * from Names join Trades
on (Names.Symbol = Trades.Symbol)
```

# Inner Join

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

Names


| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

**A copy of the smaller  
table is stored in a  
hash table like  
structure**

# Inner Join


Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |




Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |



Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |




This hash table is  
copied to each  
mappers' local disk

# Inner Join

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |




Trades

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |




Trades

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| AAPL   |      |      |     |       |     |

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |



Trades

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| MSFT   |      |      |     |       |     |

**Parts** of the larger table are distributed to each mapper

# Inner Join

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |


Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |


Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |


| Symbol |
|--------|
| GOOG   |



| Symbol |
|--------|
| AAPL   |



| Symbol |
|--------|
| MSFT   |



Mappers run on  
the **entire** Names  
and **parts** of the  
Trades table

# Inner Join


Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |



Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |



Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |



**Combine with  
those rows which  
are available on  
the mapper**



# Inner Join

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |


Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

Names


| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

| Symbol |
|--------|
| GOOG   |




| Symbol | Other Columns |
|--------|---------------|
| GOOG   |               |

| Symbol |
|--------|
| AAPL   |



| Symbol | Other Columns |
|--------|---------------|
| AAPL   |               |

| Symbol |
|--------|
| MSFT   |



| Symbol | Other Columns |
|--------|---------------|
| MSFT   |               |

The output of **all** the mappers  
forms the final output

# Inner Join

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |


Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

Names


| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

| Symbol |
|--------|
| GOOG   |




| Symbol | Other Columns |
|--------|---------------|
| GOOG   |               |

| Symbol |
|--------|
| AAPL   |



| Symbol | Other Columns |
|--------|---------------|
| AAPL   |               |

| Symbol |
|--------|
| MSFT   |



| Symbol | Other Columns |
|--------|---------------|
| MSFT   |               |

No reducer needed

The left table is the smaller table

## **Inner Joins**

Are possible as map-only joins

# Left Outer Join

**Names**

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

**Trades**

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |

All rows from the **left** table are in the result

- with a matching row
- padded with nulls

# Left Outer Join

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

Trades

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |

```
select * from Names left outer join Trades
on (Names.Symbol = Trades.Symbol)
```

# Left Outer Join

Names




| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |




The smaller table  
is copied to each  
mappers's disk

# Left Outer Join

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

| Symbol |
|--------|
| GOOG   |




Is the row **not** present in this chunk or **not** present in the entire table?

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |


| Symbol |
|--------|
| AAPL   |



Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

| Symbol |
|--------|
| MSFT   |



**No way to tell**

The left table is the smaller table

**Inner Joins**

~~**Left Outer Joins**~~

Are possible as map-only joins



# Right Outer Join

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

Trades

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |

All rows from the **right** table are in the result

- with a matching row
- padded with nulls

# Right Outer Join

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

Trades

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |


```
select * from Names right outer join Trades
on (Names.Symbol = Trades.Symbol)
```

# Right Outer Join

~~Names~~

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |


| Symbol |
|--------|
| GOOG   |



~~Names~~

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |


| Symbol |
|--------|
| AAPL   |



~~Names~~

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

| Symbol |
|--------|
| MSFT   |



The entire left table is present to check for matches

The left table is the smaller table

**Inner Joins**

~~**Left Outer Joins**~~

**Right Outer Joins**

Are possible as map-only joins

# Full Outer Join

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

Trades

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |

All rows from the **both** tables are in the result

- with a matching row
- padded with nulls

# Full Outer Join

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

| Symbol |
|--------|
| GOOG   |




Should a row be padded with nulls or does it have a match?

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

| Symbol |
|--------|
| AAPL   |



No way to tell

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

| Symbol |
|--------|
| MSFT   |



The left table is the smaller table

**Inner Joins**

~~**Left Outer Joins**~~

**Right Outer Joins**

~~**Full Outer Joins**~~

Are possible as map-only joins

# Right Table Is Small

Trades

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

What is possible as a map-only  
join now is different



The right table is the smaller table

**Inner Joins**

**Left Outer Joins**

~~**Right Outer Joins**~~

~~**Full Outer Joins**~~

Are possible as map-only joins


# Conditions for Map-side Joins



All tables, except one, are **small** enough to be held in **memory**

Tables are **bucketed on the join columns** and  
 $\text{Table1 Buckets} = N * \text{Table2 Buckets}$

# Conditions for Map-side Joins



All tables, except one, are small enough to be held in memory

Tables are **bucketed on the join columns** and  
 $\text{Table1 Buckets} = N * \text{Table2 Buckets}$

# Tables Bucketed on Join Columns

| Names  |           |
|--------|-----------|
| Symbol | Name      |
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

| Trades |      |      |     |       |     |
|--------|------|------|-----|-------|-----|
| Symbol | Open | High | Low | Close | Day |
| GOOG   |      |      |     |       |     |

**Join column = Symbol**

**Bucketed column = Symbol**

# Tables Bucketed on Join Columns

**Names**

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

**Trades**

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |

GOOG, AAPL

MSFT, FB

GOOG

AAPL

MSFT

FB

# Tables Bucketed on Join Columns

**Names**

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

**Trades**

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |

GOOG, AAPL

MSFT, FB

GOOG

AAPL

MSFT

FB

# Tables Bucketed on Join Columns

**Names**

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

**Trades**

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |

GOOG, AAPL

MSFT, FB

GOOG

AAPL

MSFT

FB

# Bucket Map Join

Names

| Symbol | Name      |
|--------|-----------|
| GOOG   | Google    |
| AAPL   | Apple     |
| MSFT   | Microsoft |

Trades

| Symbol | Open | High | Low | Close | Day |
|--------|------|------|-----|-------|-----|
| GOOG   |      |      |     |       |     |

GOOG, AAPL

MSFT, FB

GOOG

AAPL

MSFT

FB



```
set hive.optimize.bucketmapjoin = true;
```

---

## Enable Bucket Map Join

**This is not enabled by default in Hive**

Bucket map joins are more efficient when tables are...

**Bucketed** on the join column

**Sorted** on the join column

Have the **same number** of buckets

Rows can be joined using  
**merge-sort**

Can only be used for **equi-joins**

```
set hive.input.format =  
org.apache.hadoop.hive.q1.io.BucketizedHiveInputFormat;
```

```
set hive.optimize.bucketmapjoin = true;
```

```
set hive.optimize.bucketmapjoin.sortedmerge = true;
```

---

## Enable Bucket Map Join with Sort-merge

**Format of the file that is read in from disk into Hive**

```
set hive.input.format =  
org.apache.hadoop.hive.q1.io.BucketizedHiveInputFormat;
```

```
set hive.optimize.bucketmapjoin = true;
```

```
set hive.optimize.bucketmapjoin.sortedmerge = true;
```

---

## Enable Bucket Map Join with Sort-merge

**Format of the file that is read in from disk into Hive**

```
set hive.input.format =  
org.apache.hadoop.hive.q1.io.BucketizedHiveInputFormat;
```

```
set hive.optimize.bucketmapjoin = true;
```

```
set hive.optimize.bucketmapjoin.sortedmerge = true;
```

---

## Enable Bucket Map Join with Sort-merge

**Format of the file that is read in from disk into Hive**

```
select /*+ mapjoin(trades) */ names.symbol, high, day  
from names join trades on  
(names.symbol = trades.symbol)
```

---

## Explicitly Perform a Map-join

**Read the table specified completely on the mapper node**

# Summary

**A deep understanding of how joins work in Hive**

**Faster joins on large tables**

**Optimized semi-joins in place of IN/ EXISTS subqueries**

**Understood under what conditions joins are map-only and how to optimize them**